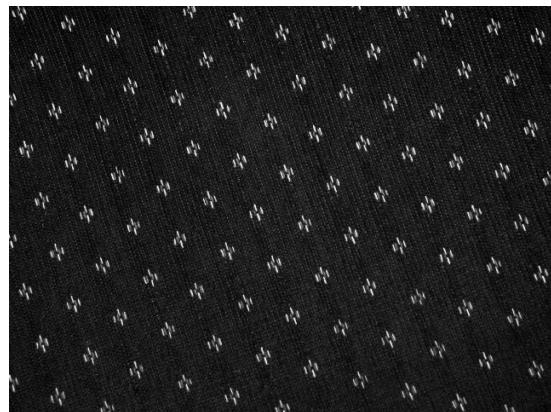
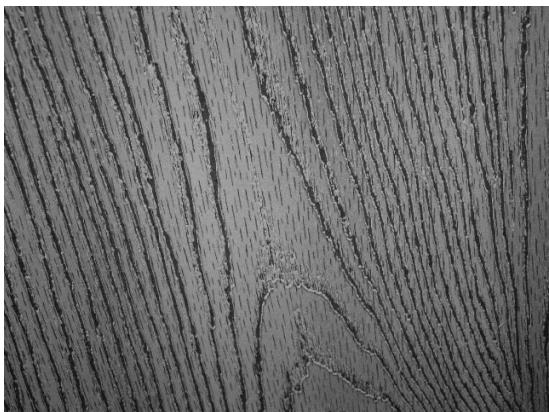


# Application: Representing Texture

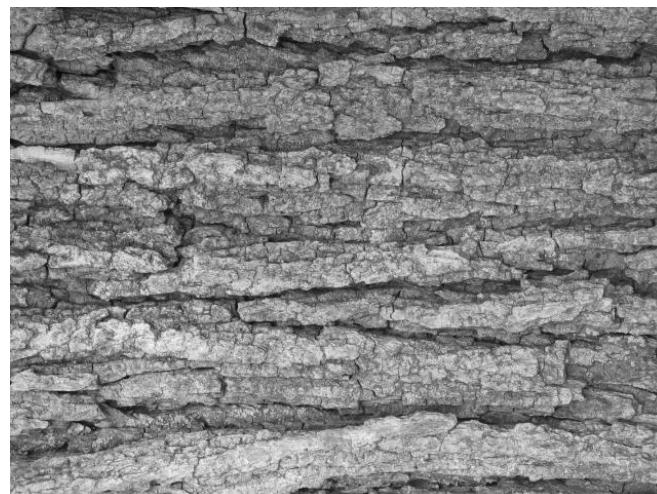


Source: Forsyth

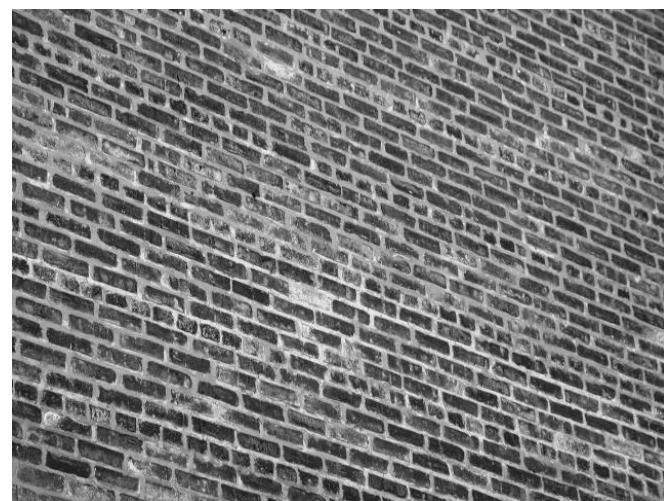
# Texture and Material



# Texture and Orientation



# Texture and Scale



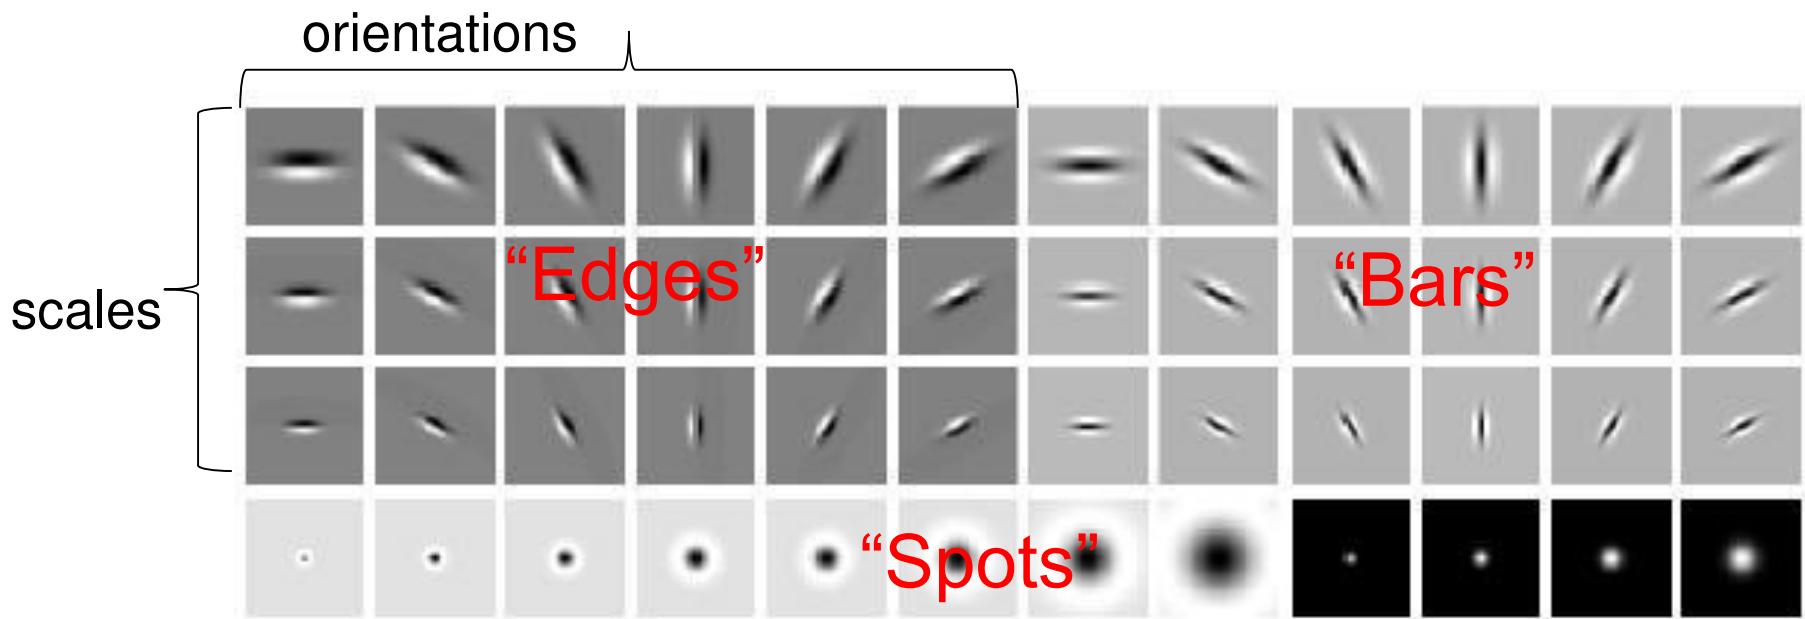
# What is texture?

Regular or stochastic patterns caused by bumps,  
grooves, and/or markings

# How can we represent texture?

- Compute responses of blobs and edges at various orientations and scales

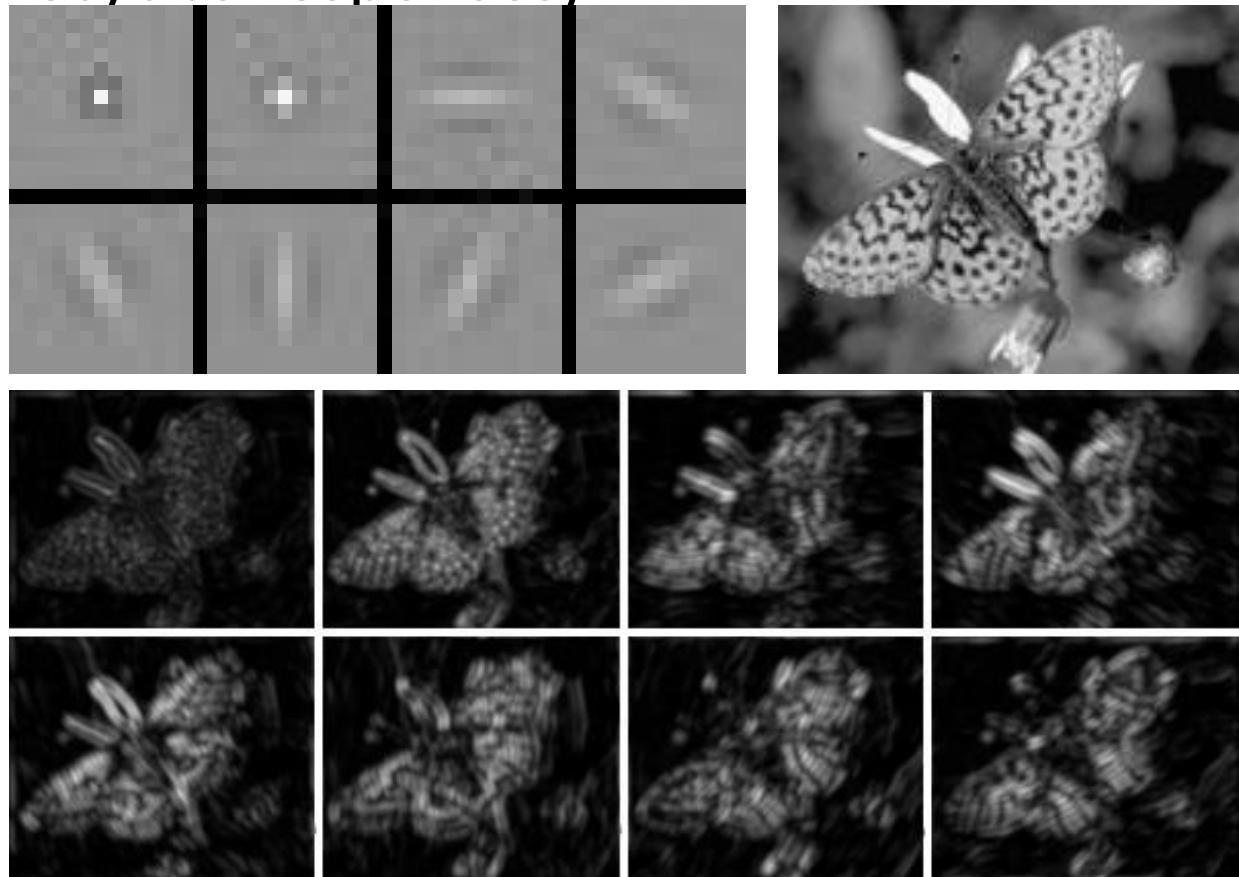
# Overcomplete representation: filter banks



Code for filter banks: [www.robots.ox.ac.uk/~vgg/research/texclass/filters.html](http://www.robots.ox.ac.uk/~vgg/research/texclass/filters.html)

# Filter banks

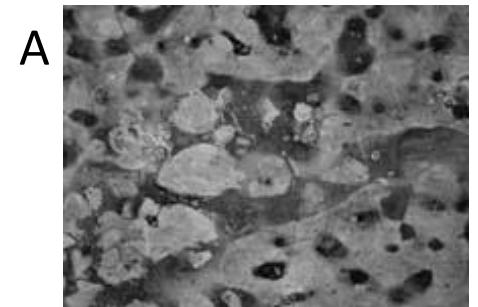
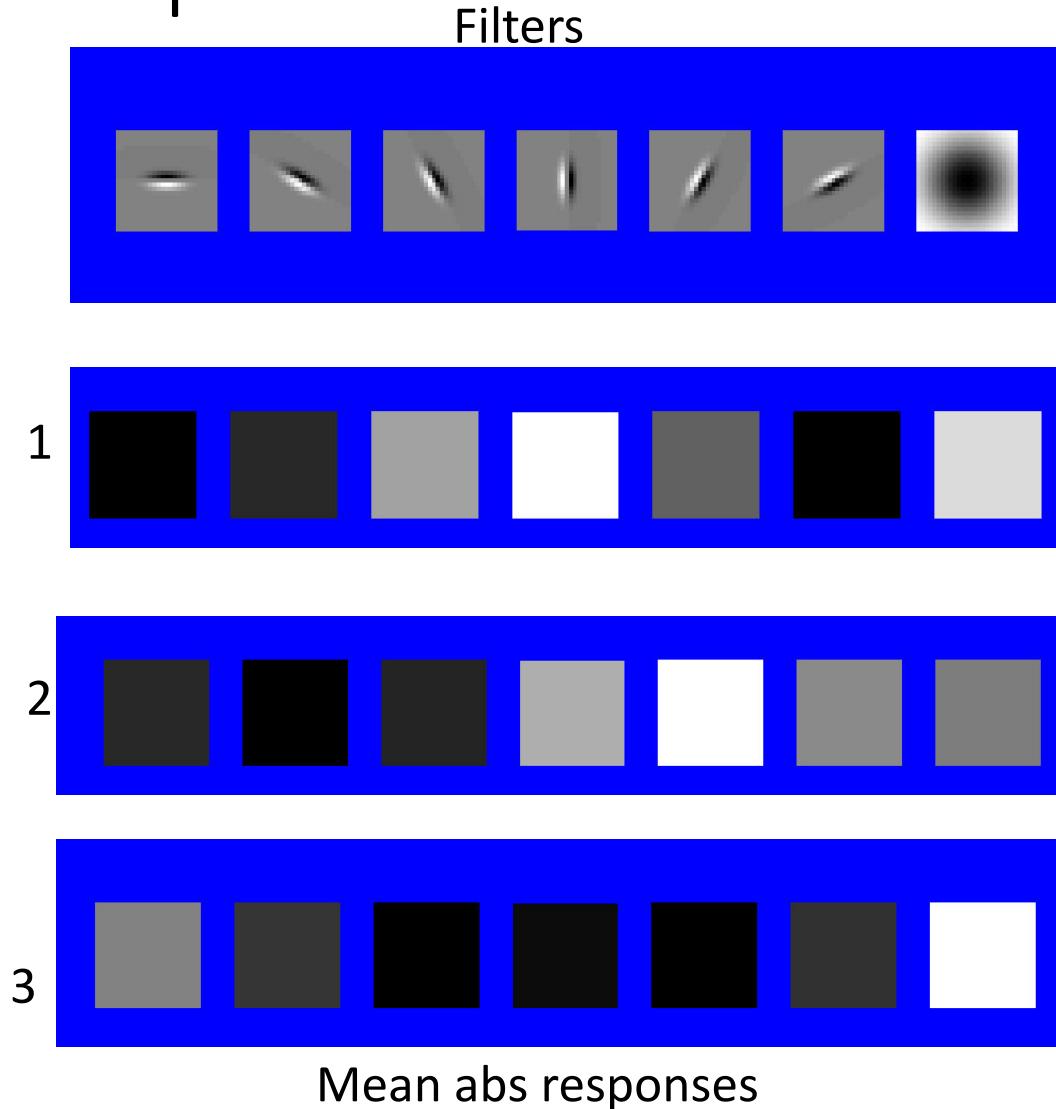
- Process image with each filter and keep responses (or squared/abs responses)



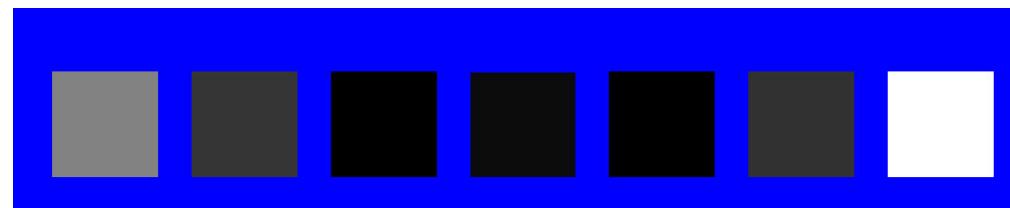
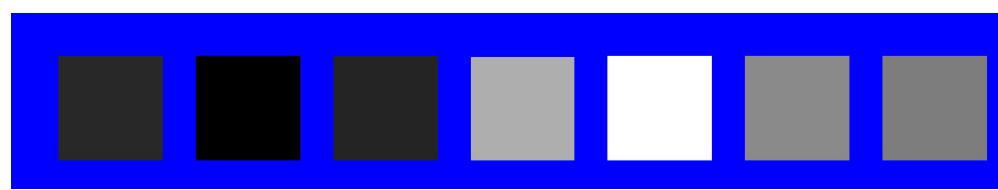
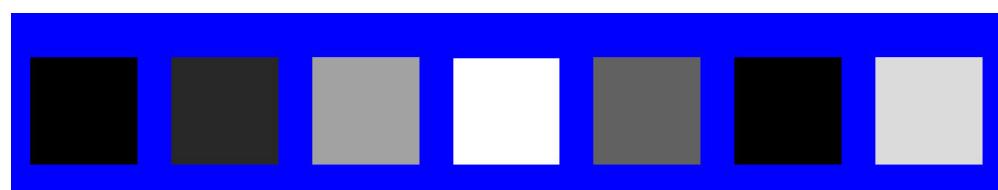
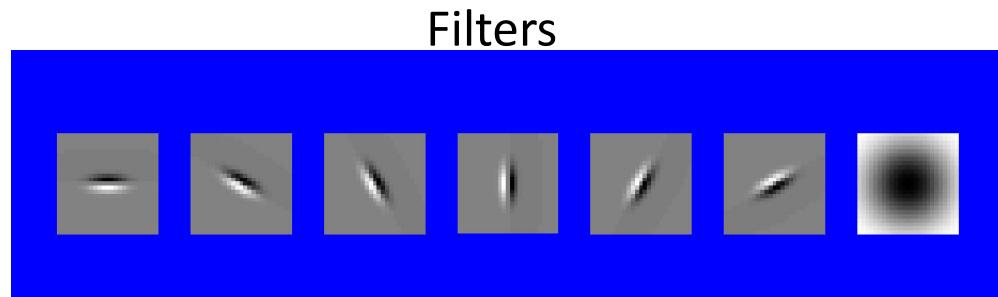
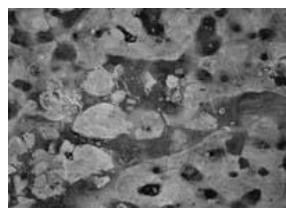
# How can we represent texture?

- Measure responses of blobs and edges at various orientations and scales
- Idea 1: Record simple statistics (e.g., mean, std.) of absolute filter responses

# Can you match the texture to the response?



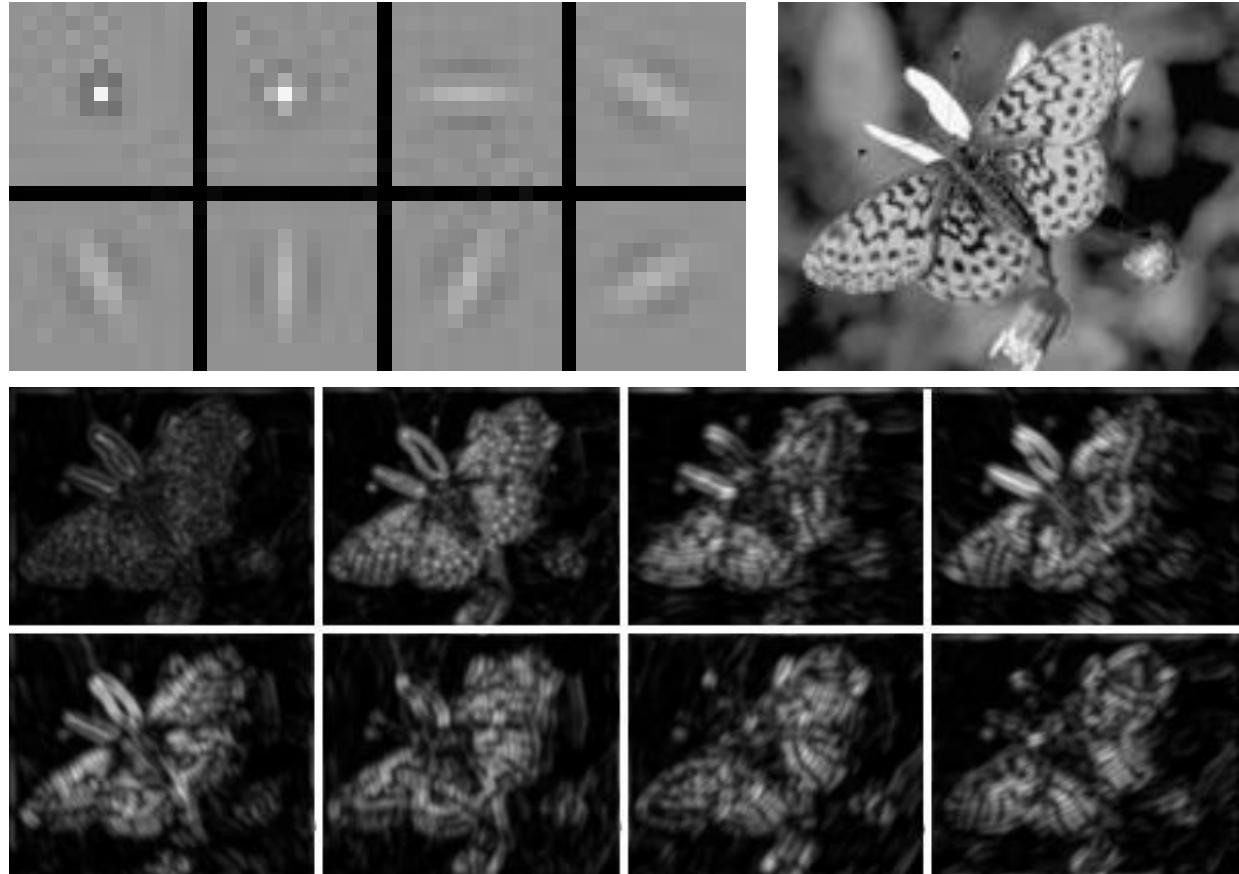
# Representing texture by mean abs response



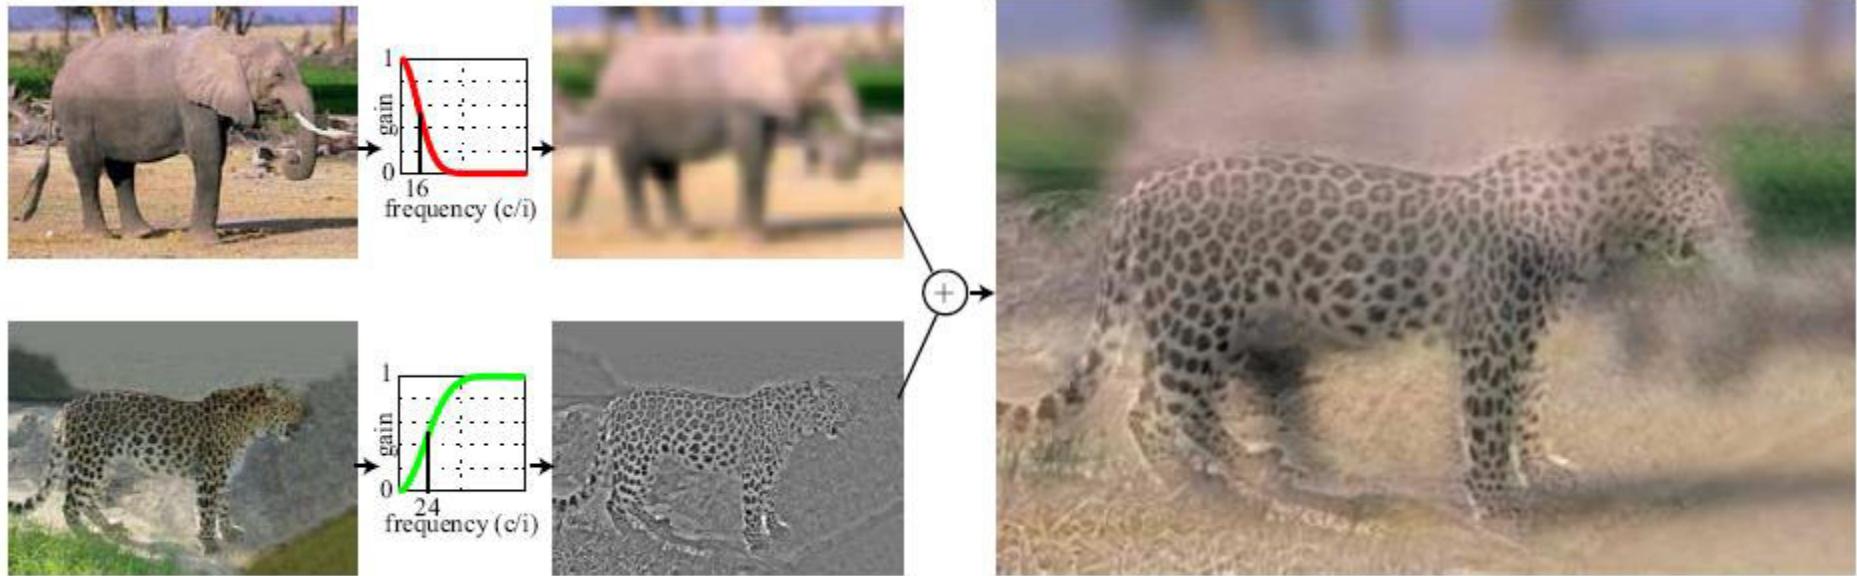
Mean abs responses

# Representing texture

- Idea 2: take vectors of filter responses at each pixel and cluster them, then take histograms (more on this in coming weeks)

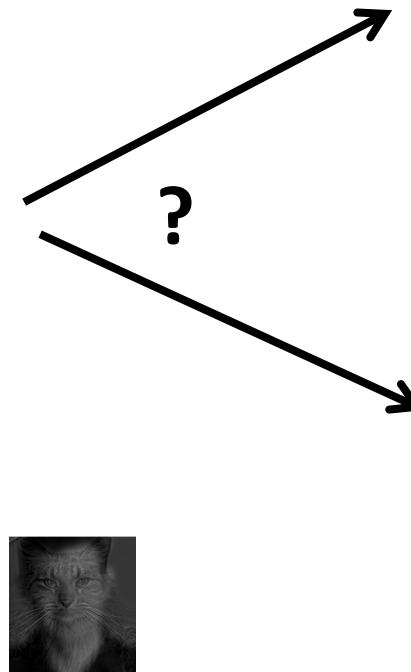


# Hybrid Images



- A. Oliva, A. Torralba, P.G. Schyns,  
“Hybrid Images,” SIGGRAPH 2006

# Why do we get different, distance-dependent interpretations of hybrid images?



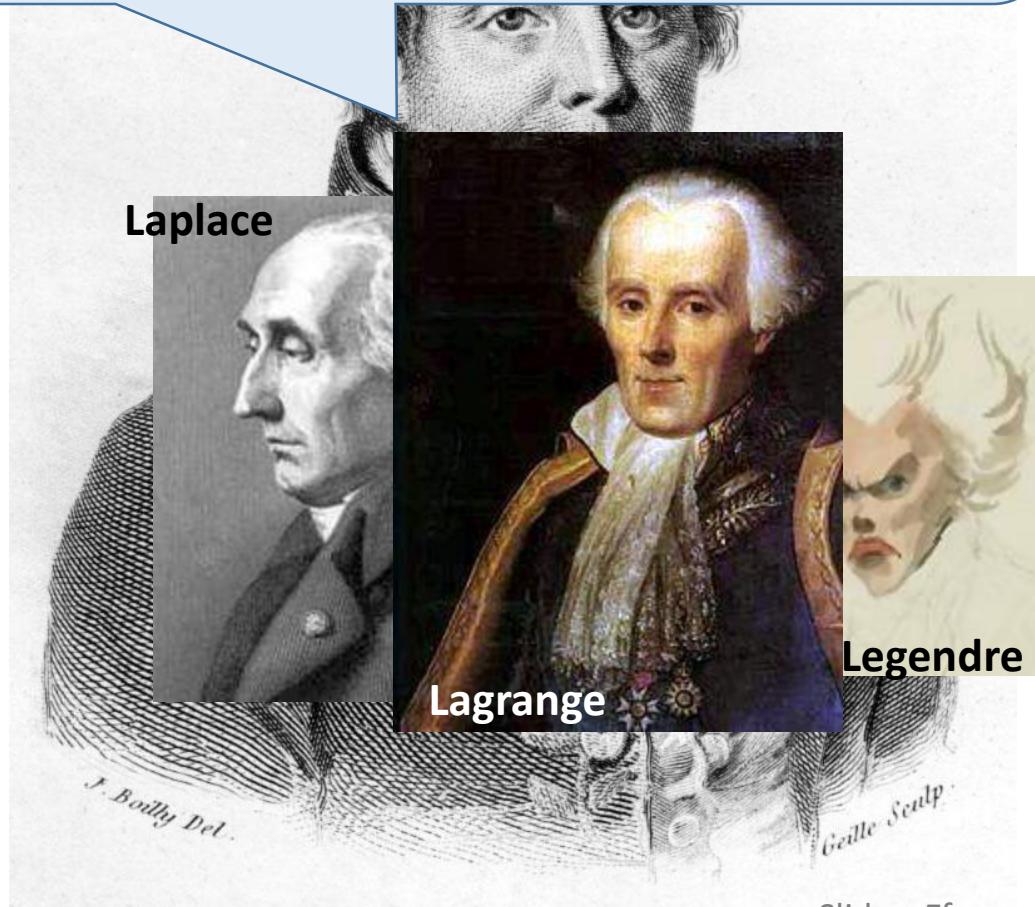
# Jean Baptiste Joseph Fourier (1768-1830)

had crazy idea (1807):

*Any univariate function can be rewritten as a weighted sum of sines and cosines of different frequencies.*

- Don't believe it?
  - Neither did Lagrange, Laplace, Poisson and other big wigs
  - Not translated into English until 1878!
- But it's (mostly) true!
  - called Fourier Series
  - there are some subtle restrictions

*...the manner in which the author arrives at these equations is not exempt of difficulties and...his analysis to integrate them still leaves something to be desired on the score of generality and even rigour.*



# Fourier, Joseph (1768-1830)



French mathematician who discovered that any periodic motion can be written as a superposition of sinusoidal and cosinusoidal vibrations. He developed a mathematical theory of [heat](#) in *Théorie Analytique de la Chaleur* (*Analytic Theory of Heat*), (1822), discussing it in terms of differential equations.

Fourier was a friend and advisor of Napoleon. Fourier believed that his health would be improved by wrapping himself up in blankets, and in this state he tripped down the stairs in his house and killed himself. The paper of [Galois](#) which he had taken home to read shortly before his death was never recovered.

**SEE ALSO:** [Galois](#)

*Additional biographies:* [MacTutor](#) (St. Andrews), [Bonn](#)

© 1996-2007 Eric W. Weisstein

How would math have changed if the Slanket or Snuggie had been invented?

Slide credit: James Hays

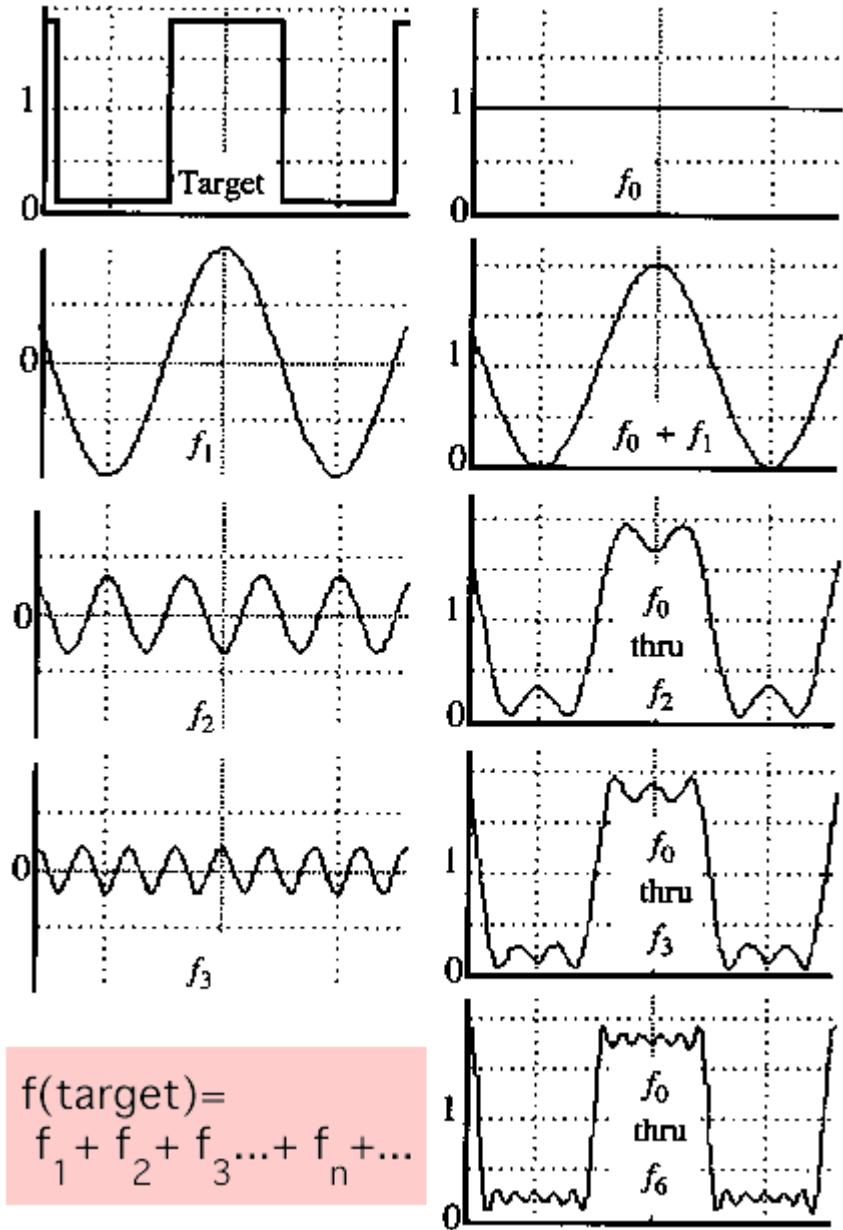


# A sum of sines

Our building block:

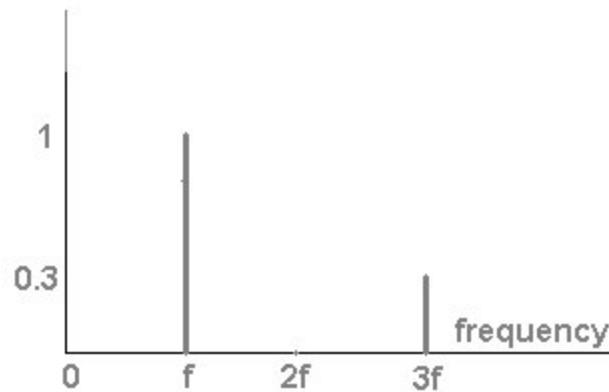
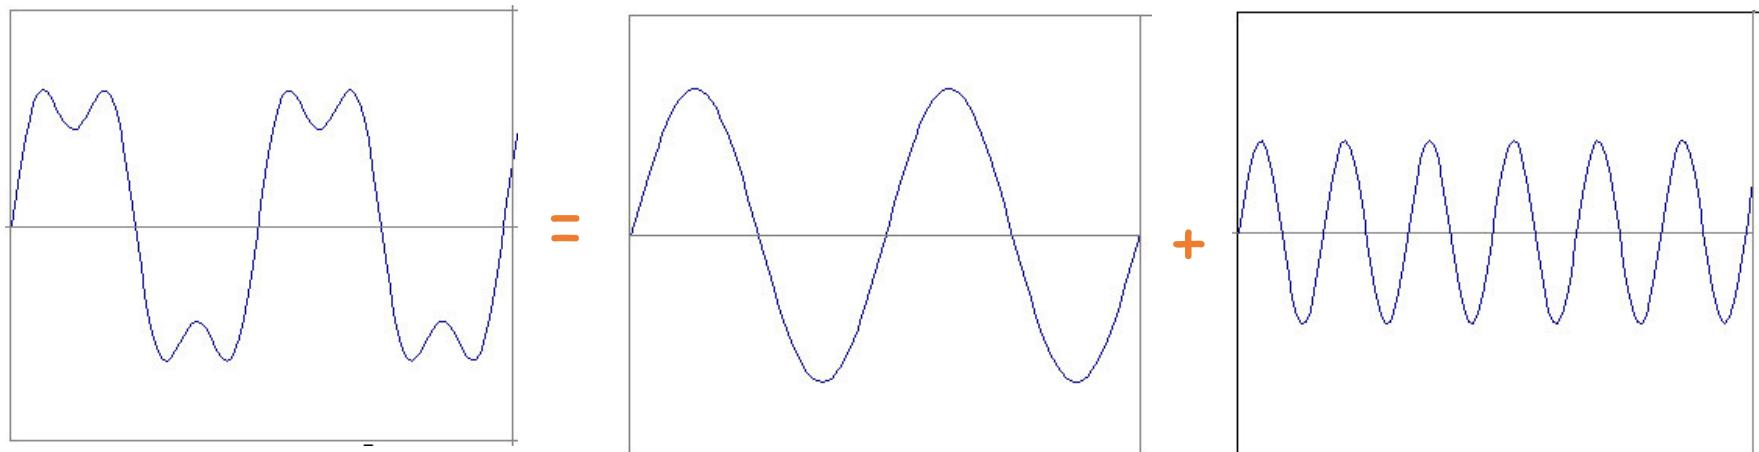
$$A \sin(\omega x + \phi)$$

Add enough of them to get any signal  $f(x)$  you want!

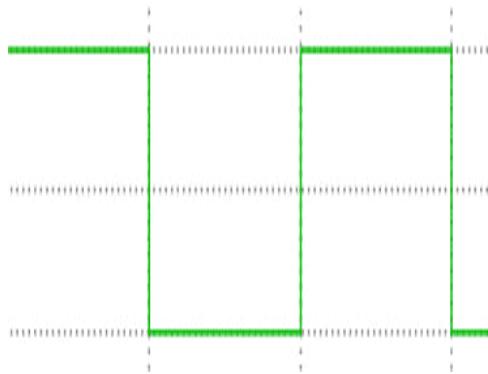


# Frequency Spectra

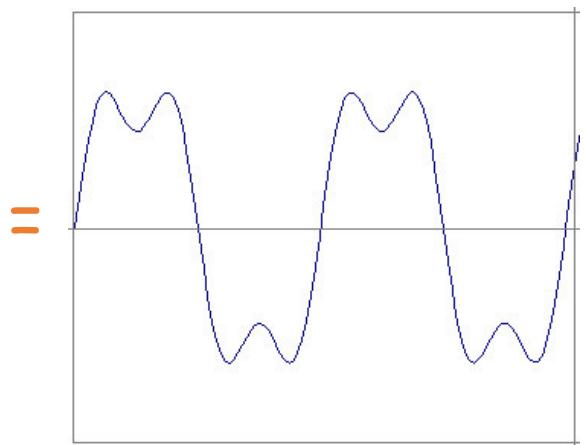
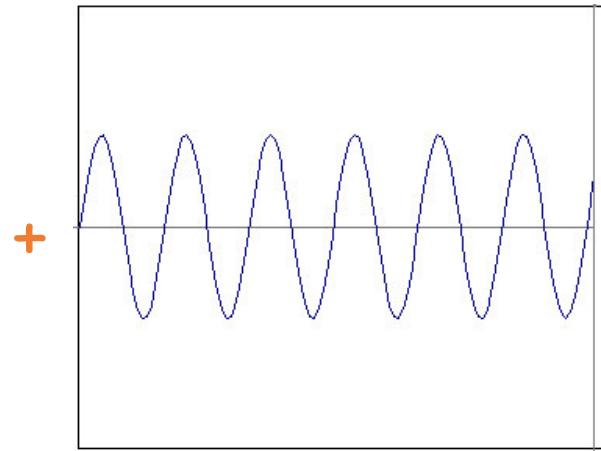
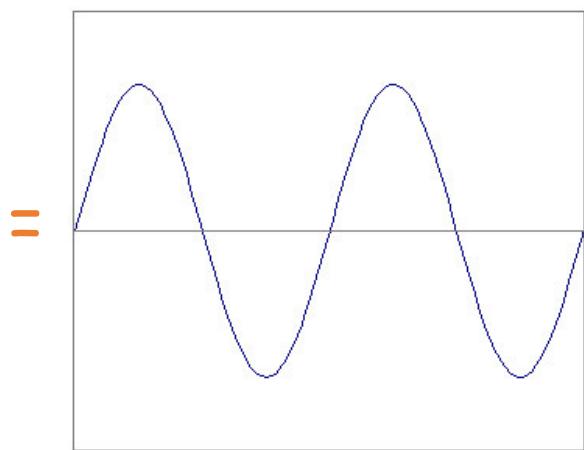
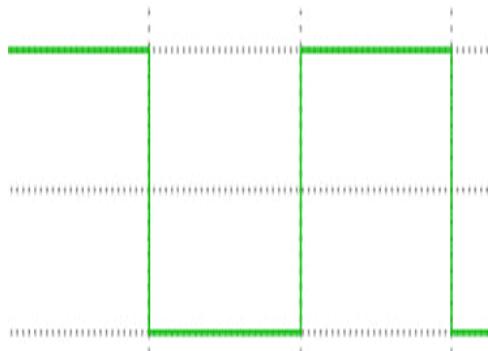
- example :  $g(t) = \sin(2\pi f t) + (1/3)\sin(2\pi(3f)t)$



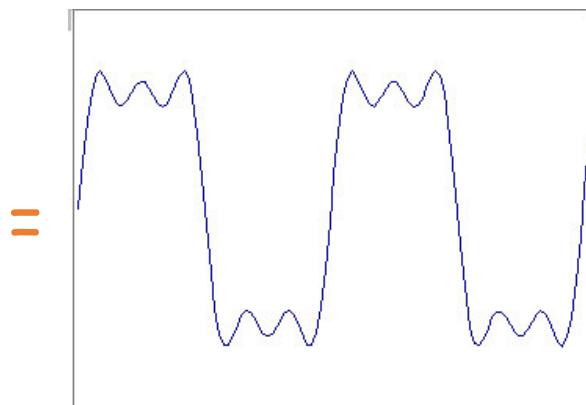
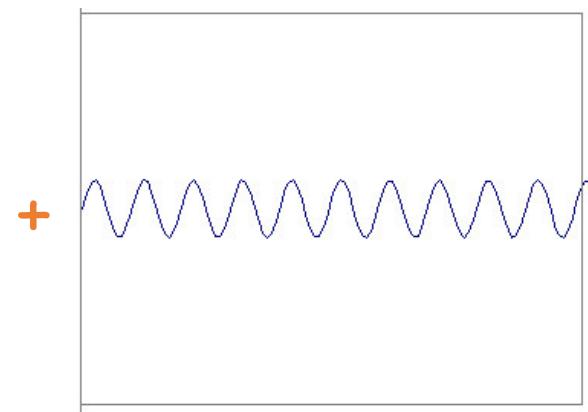
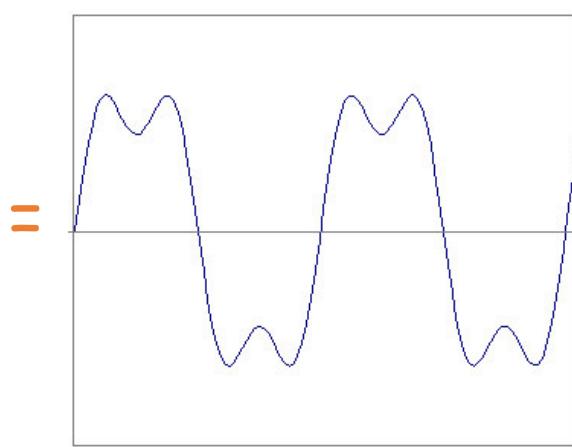
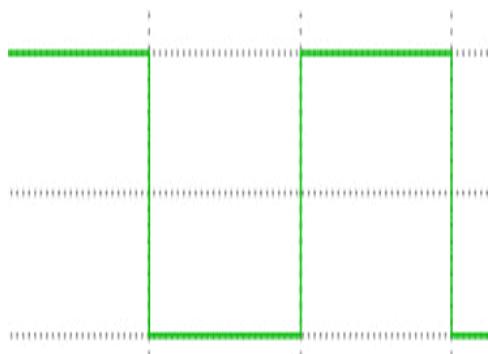
# Frequency Spectra



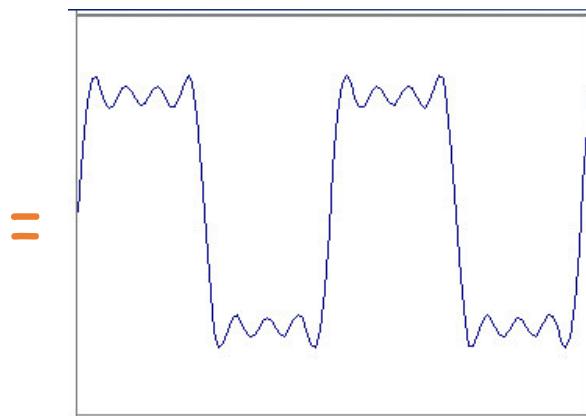
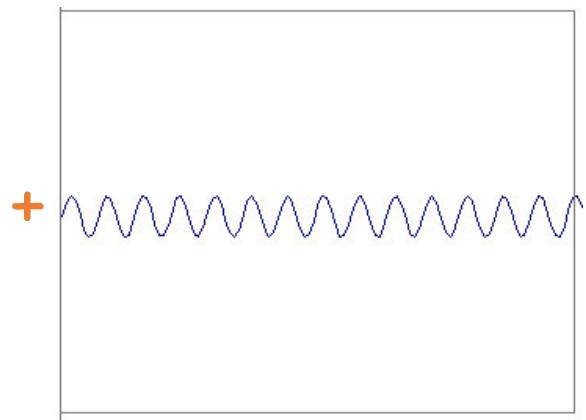
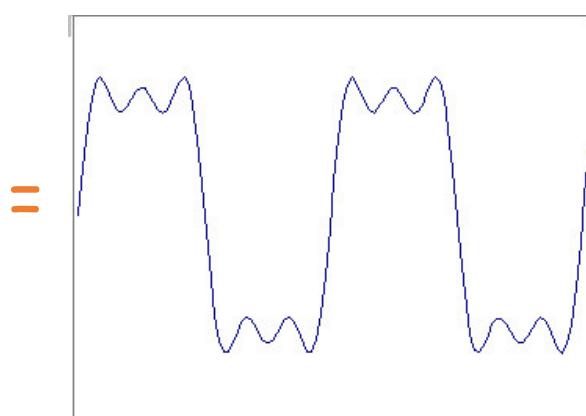
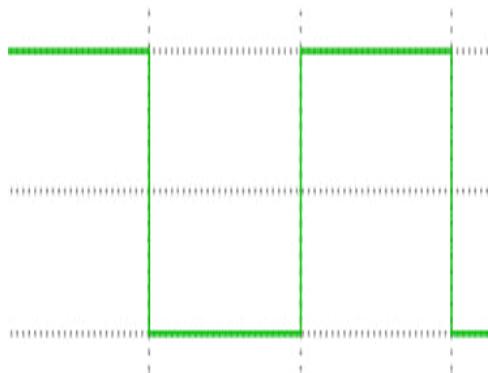
# Frequency Spectra



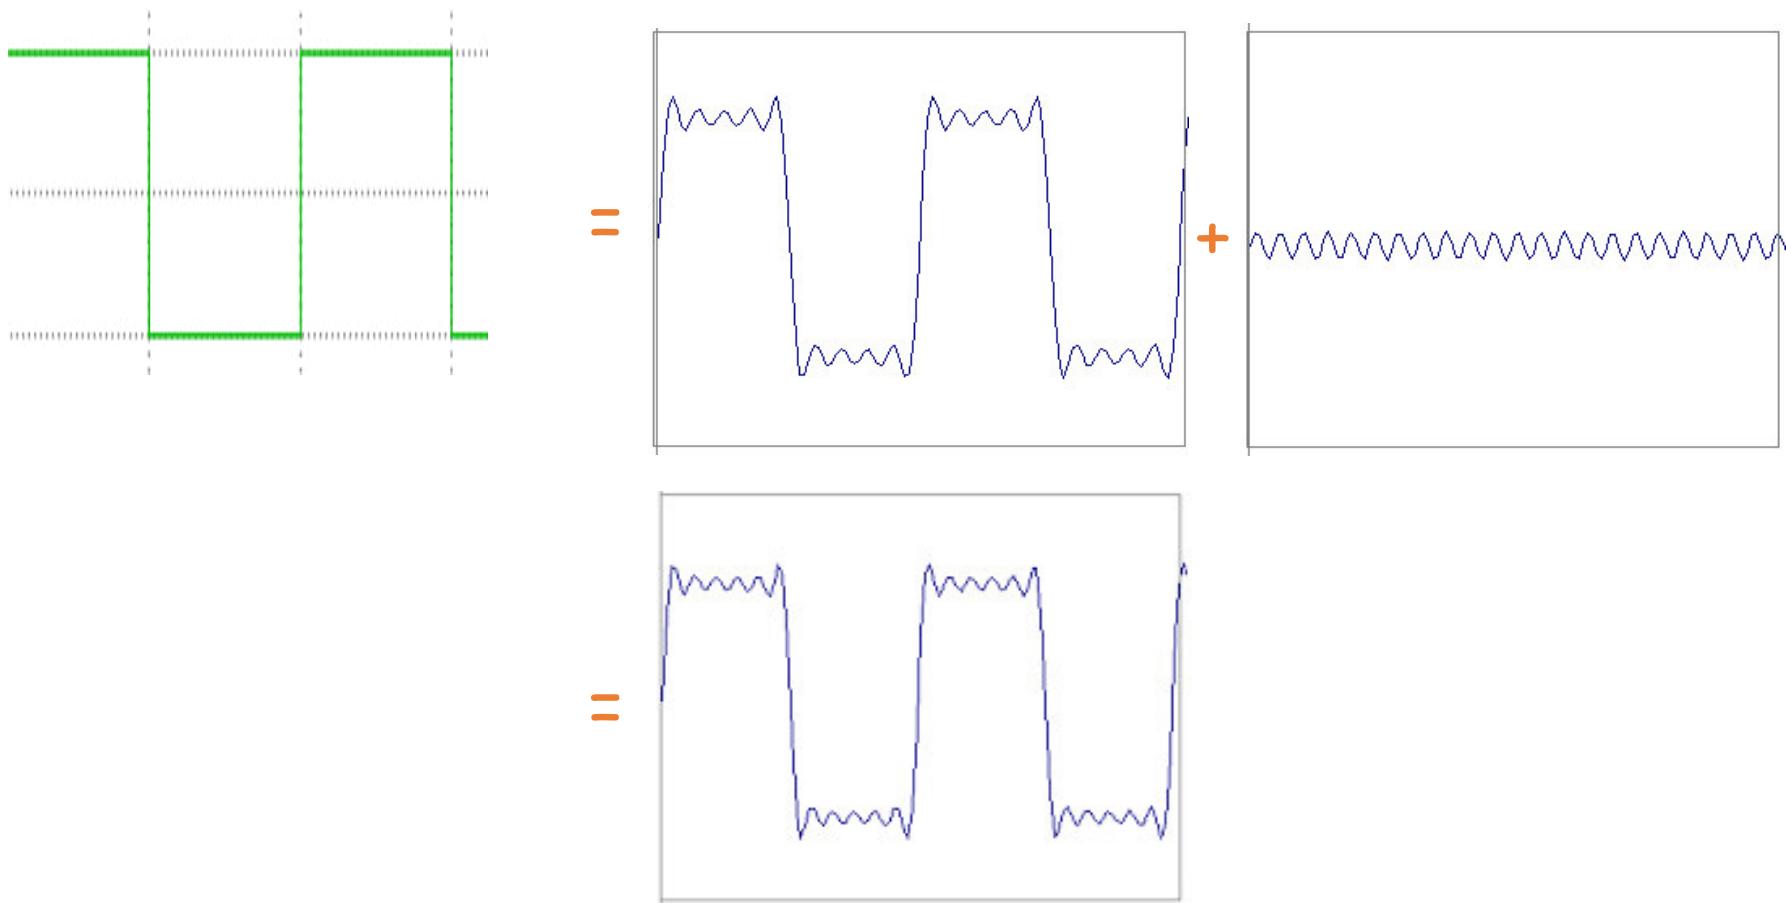
# Frequency Spectra



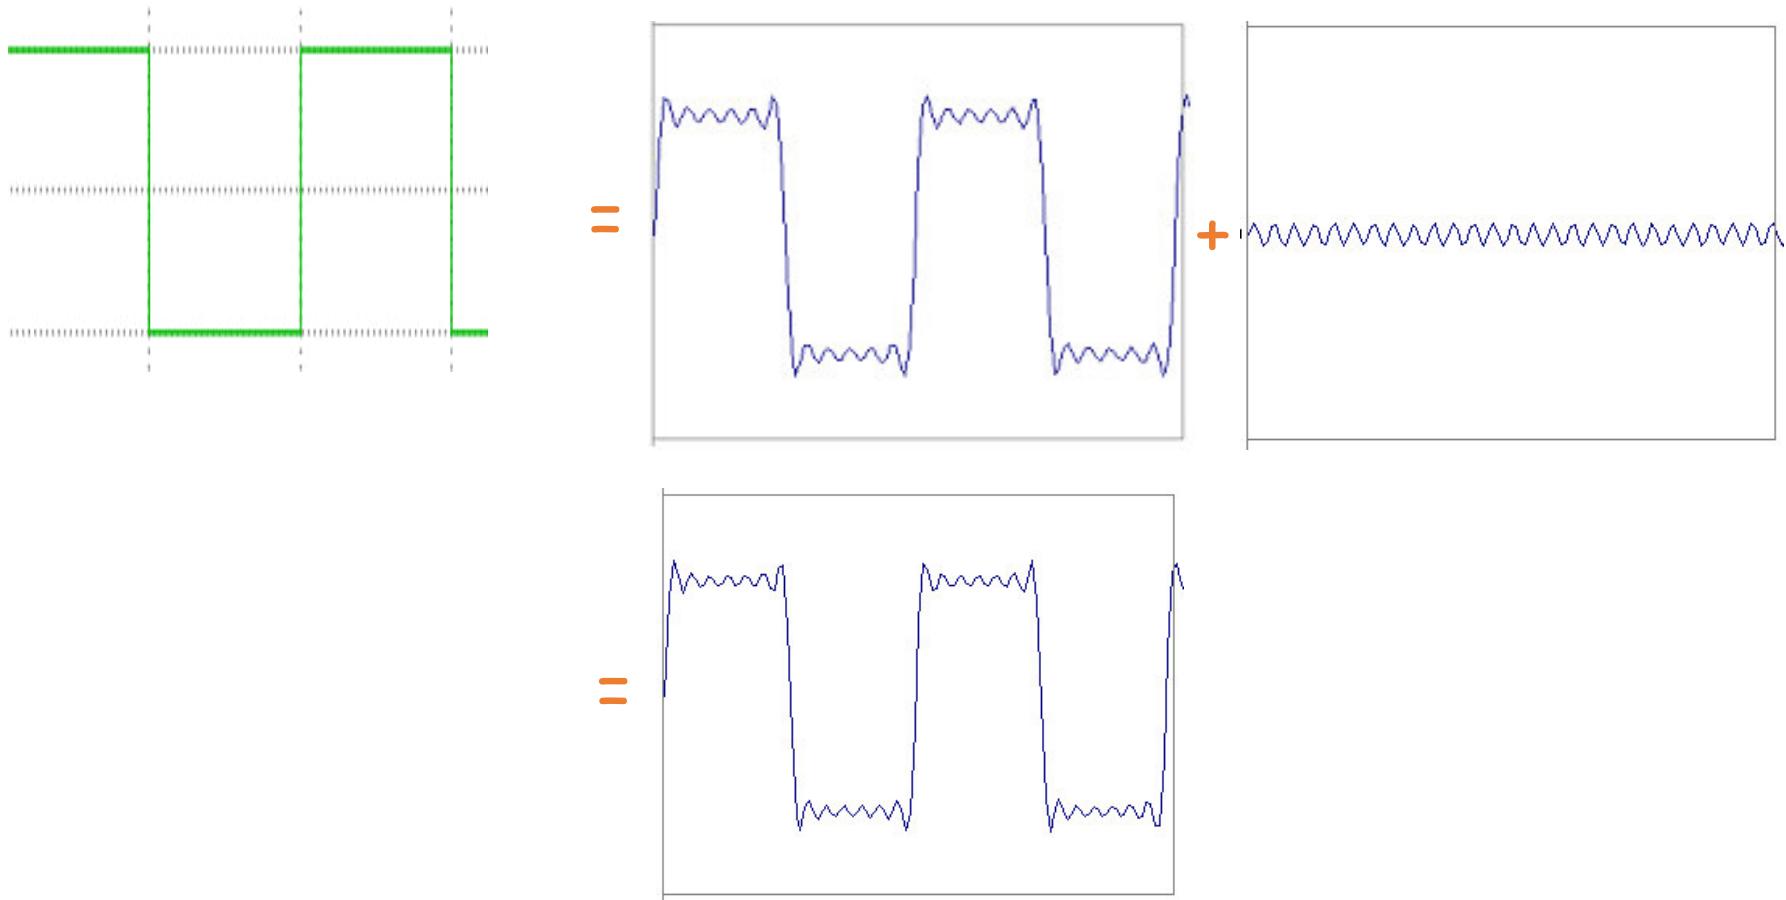
# Frequency Spectra



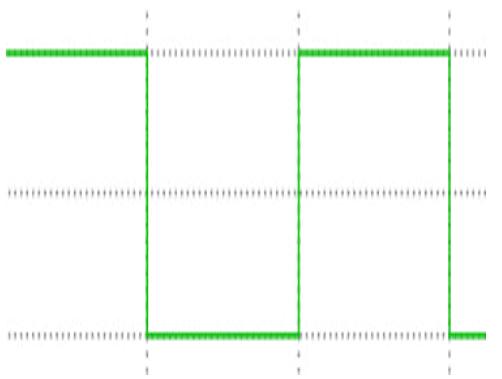
# Frequency Spectra



# Frequency Spectra

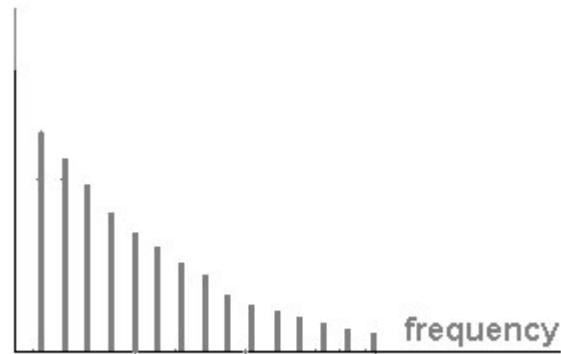


# Frequency Spectra



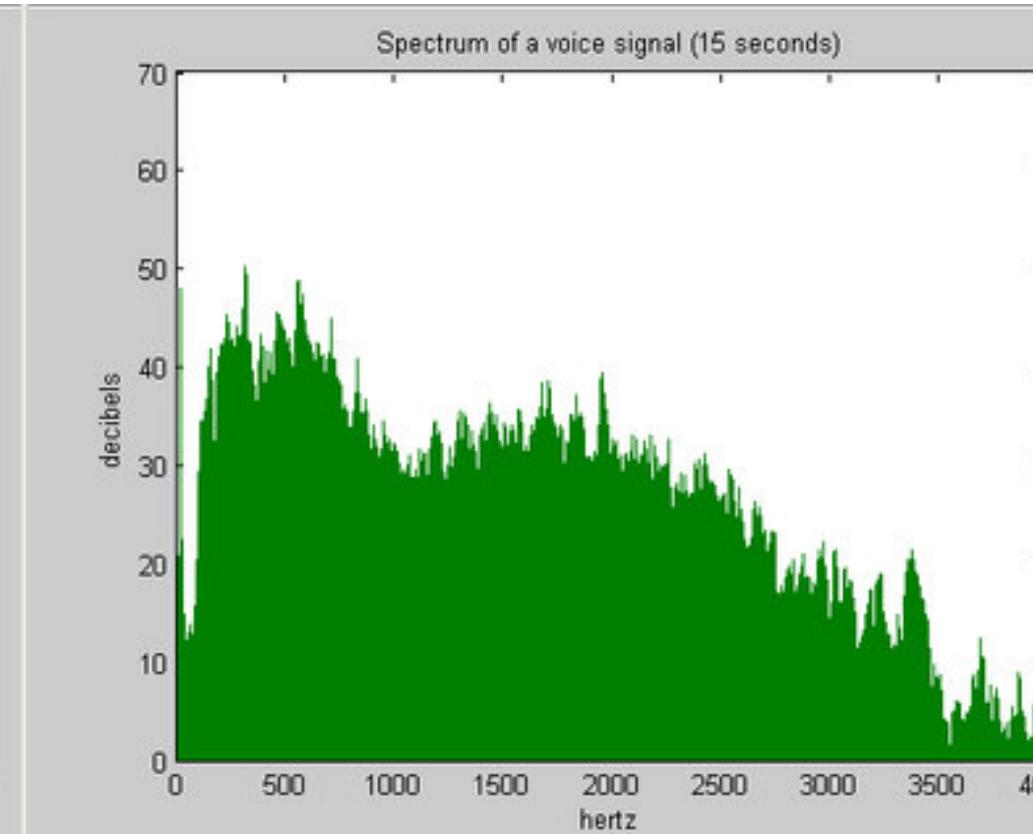
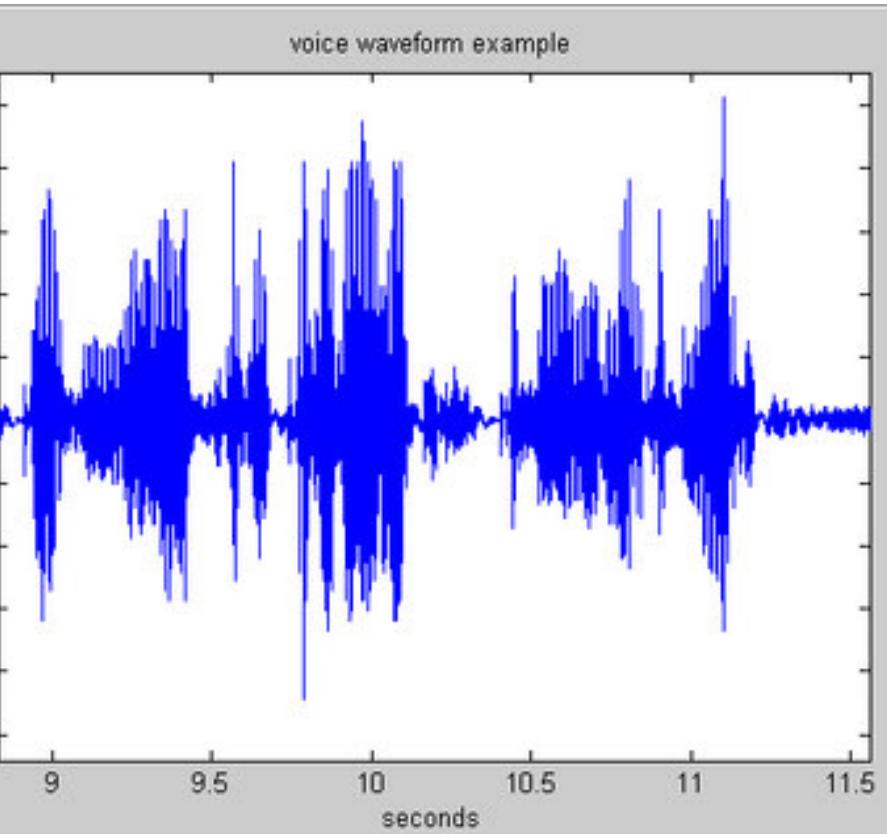
=

$$A \sum_{k=1}^{\infty} \frac{1}{k} \sin(2\pi kt)$$



# Example: Music

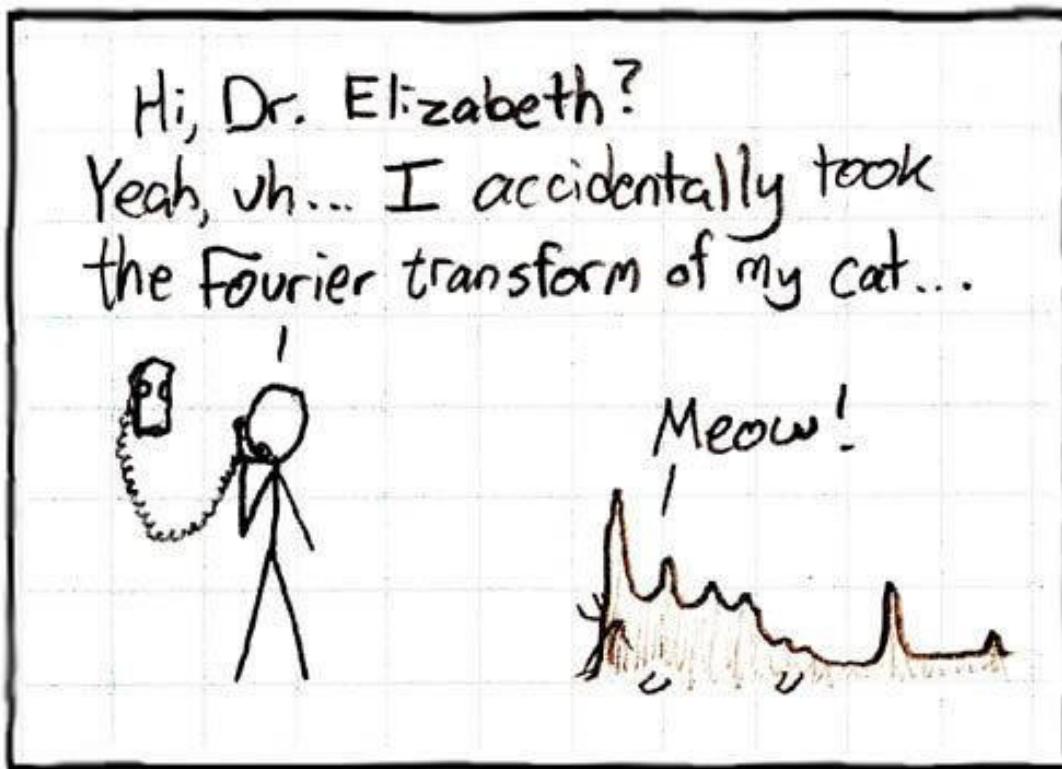
- We think of music in terms of frequencies at different magnitudes



# Other signals

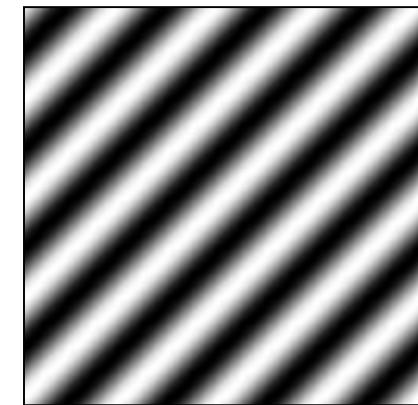
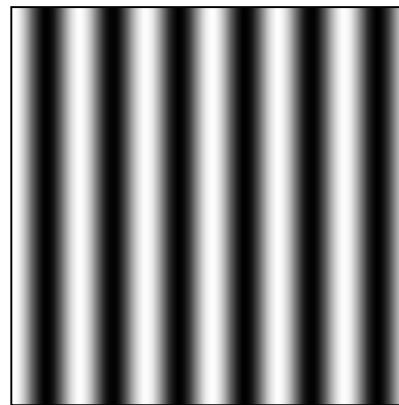
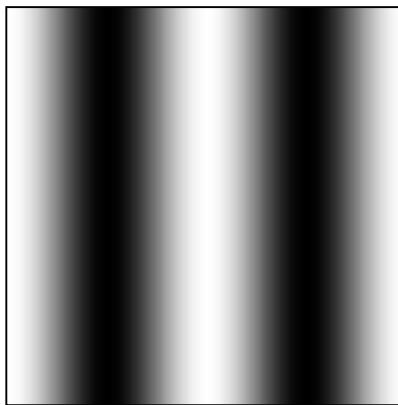
- We can also think of all kinds of other signals the same way

Cats(?)

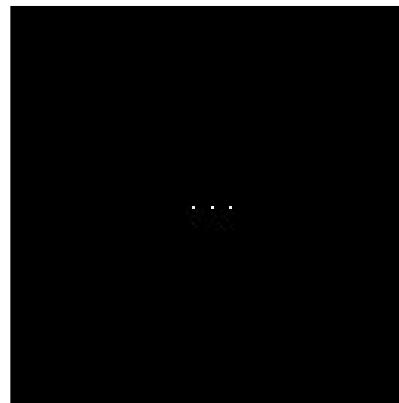
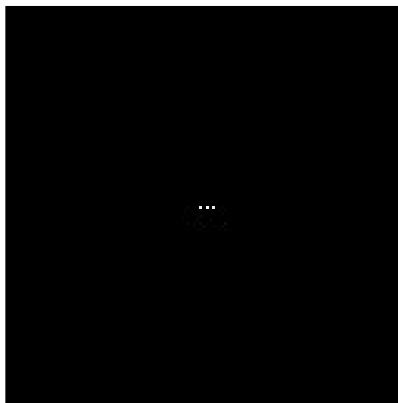


# Fourier analysis in images

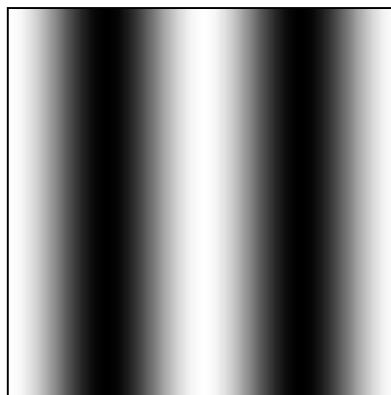
Intensity  
Image



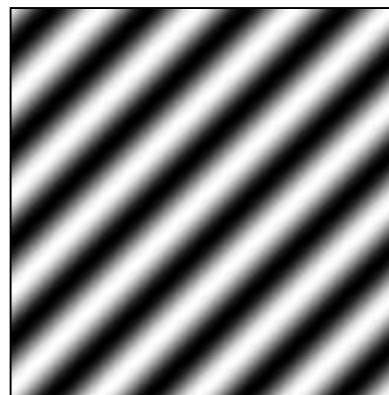
Fourier  
Image



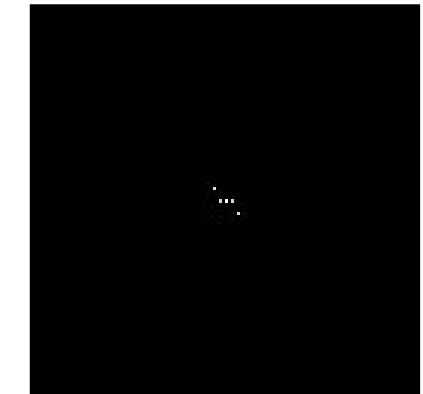
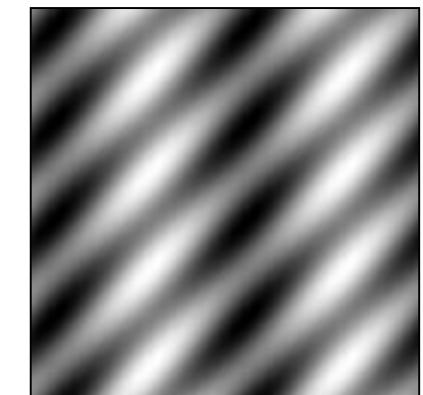
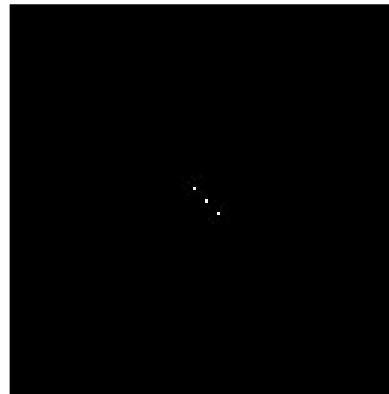
# Signals can be composed



+



=



<http://sharp.bu.edu/~slehar/fourier/fourier.html#filtering>  
More: <http://www.cs.unm.edu/~brayer/vision/fourier.html>

# Fourier Transform

- Fourier transform stores the magnitude and phase at each frequency
  - Magnitude encodes how much signal there is at a particular frequency
  - Phase encodes spatial information (indirectly)
  - For mathematical convenience, this is often notated in terms of real and complex numbers

Amplitude:  $A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$       Phase:  $\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$

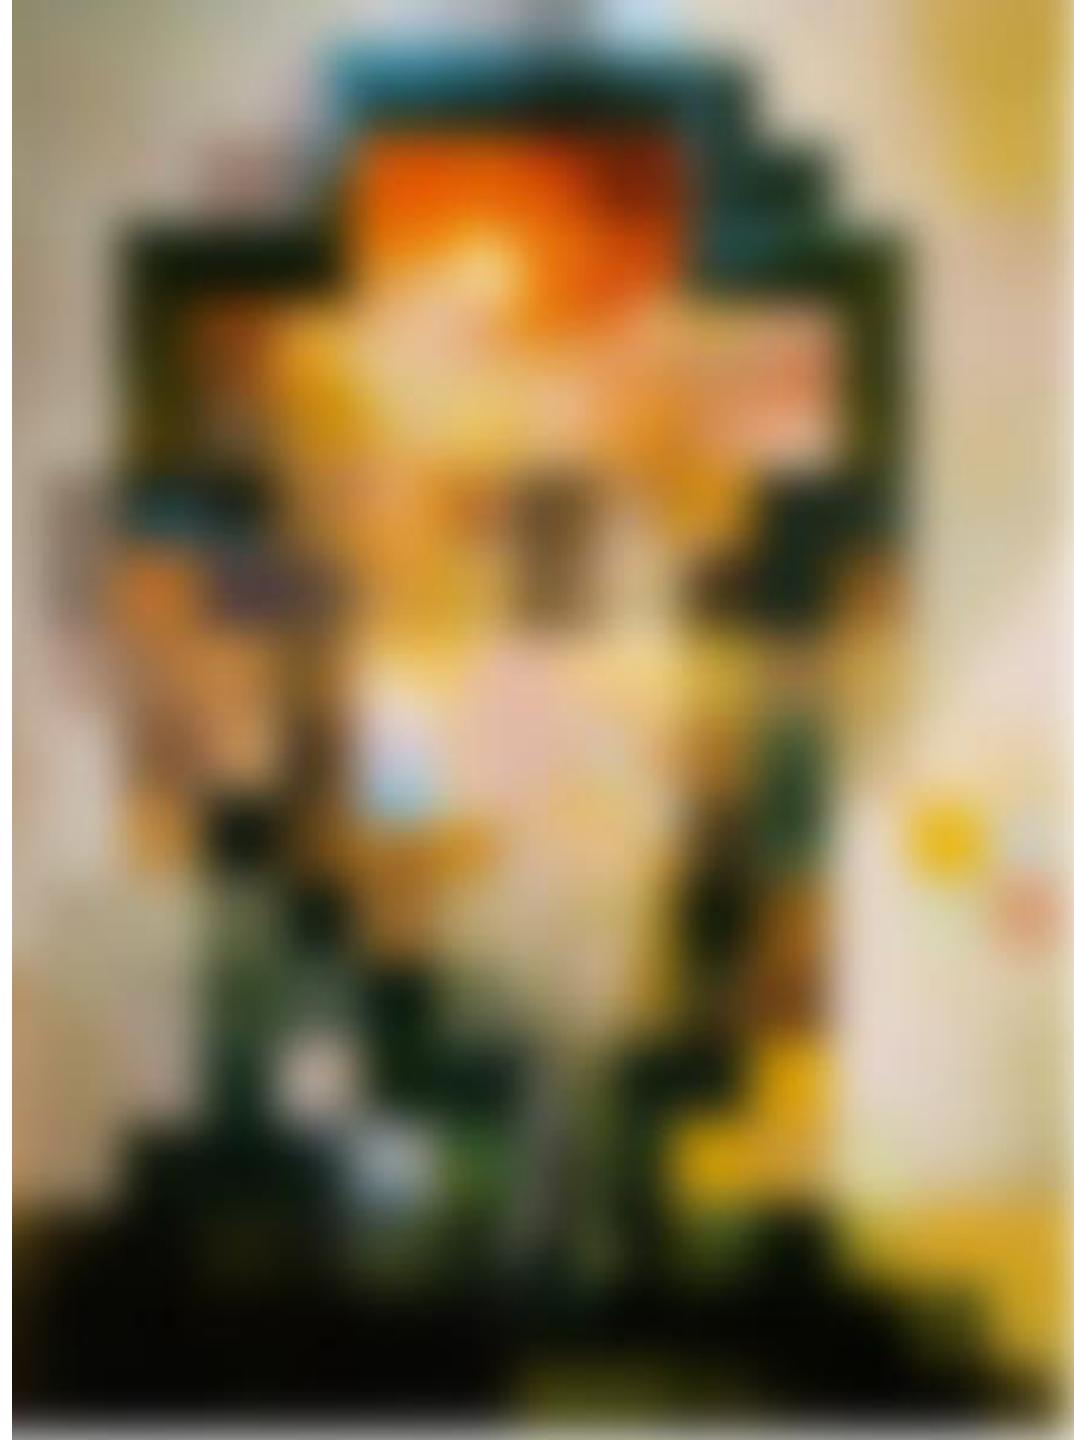
Euler's formula:  $e^{inx} = \cos(nx) + i \sin(nx)$

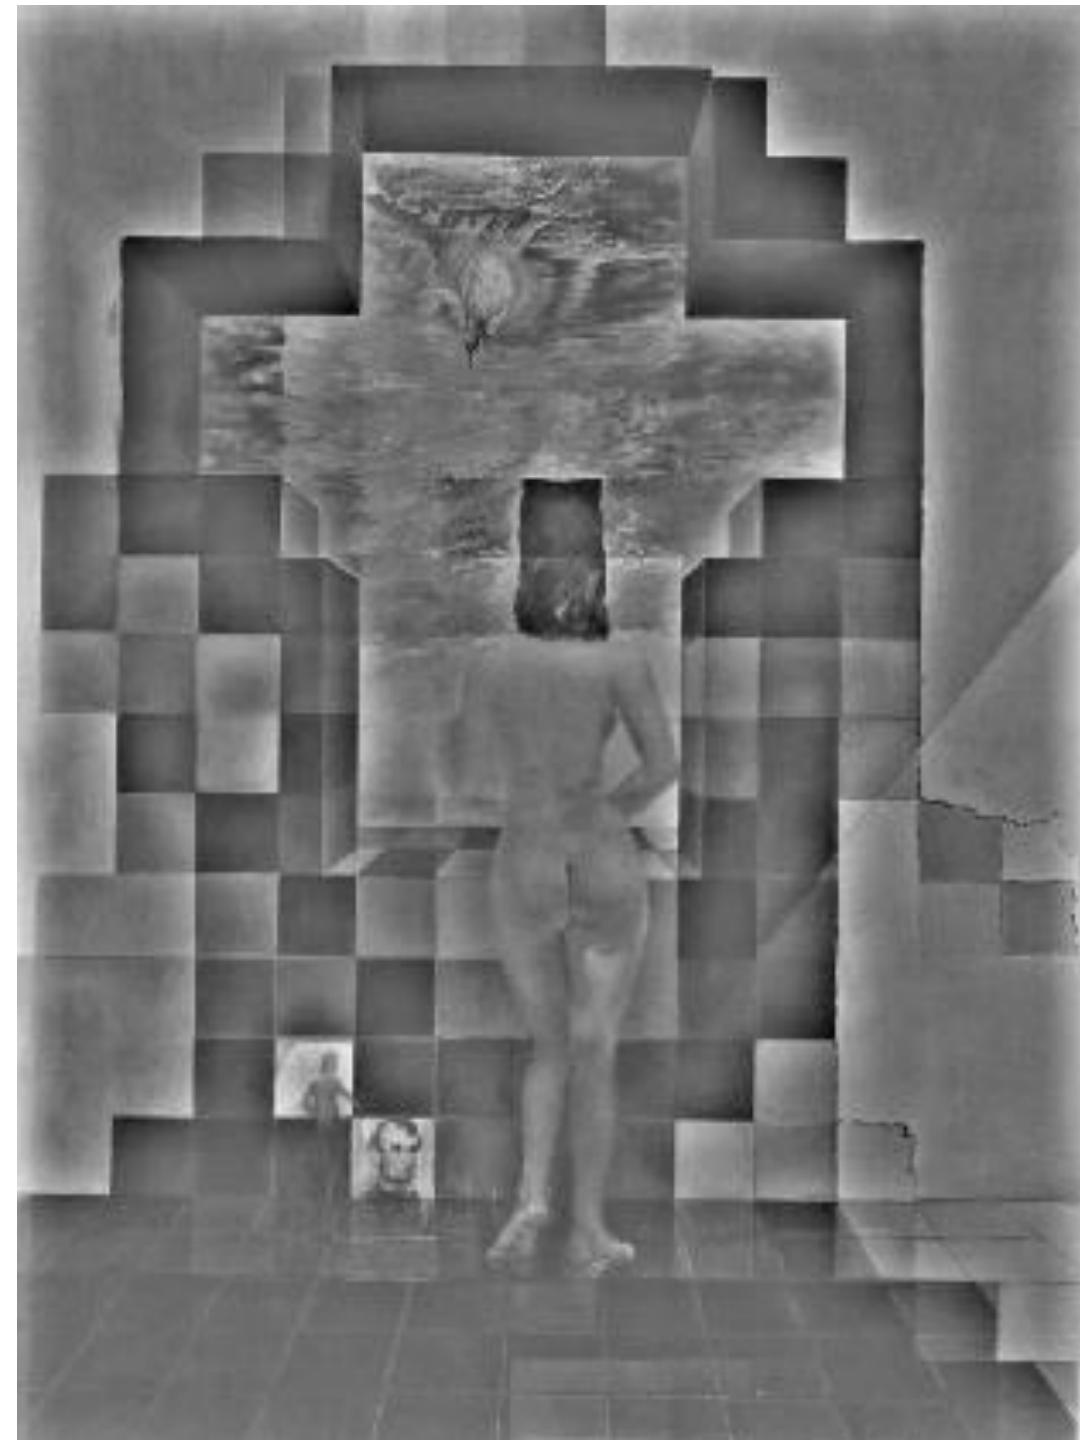
**Salvador Dali invented Hybrid Images?**

**Salvador Dali**

*“Gala Contemplating the Mediterranean Sea,  
which at 30 meters becomes the portrait  
of Abraham Lincoln”, 1976*

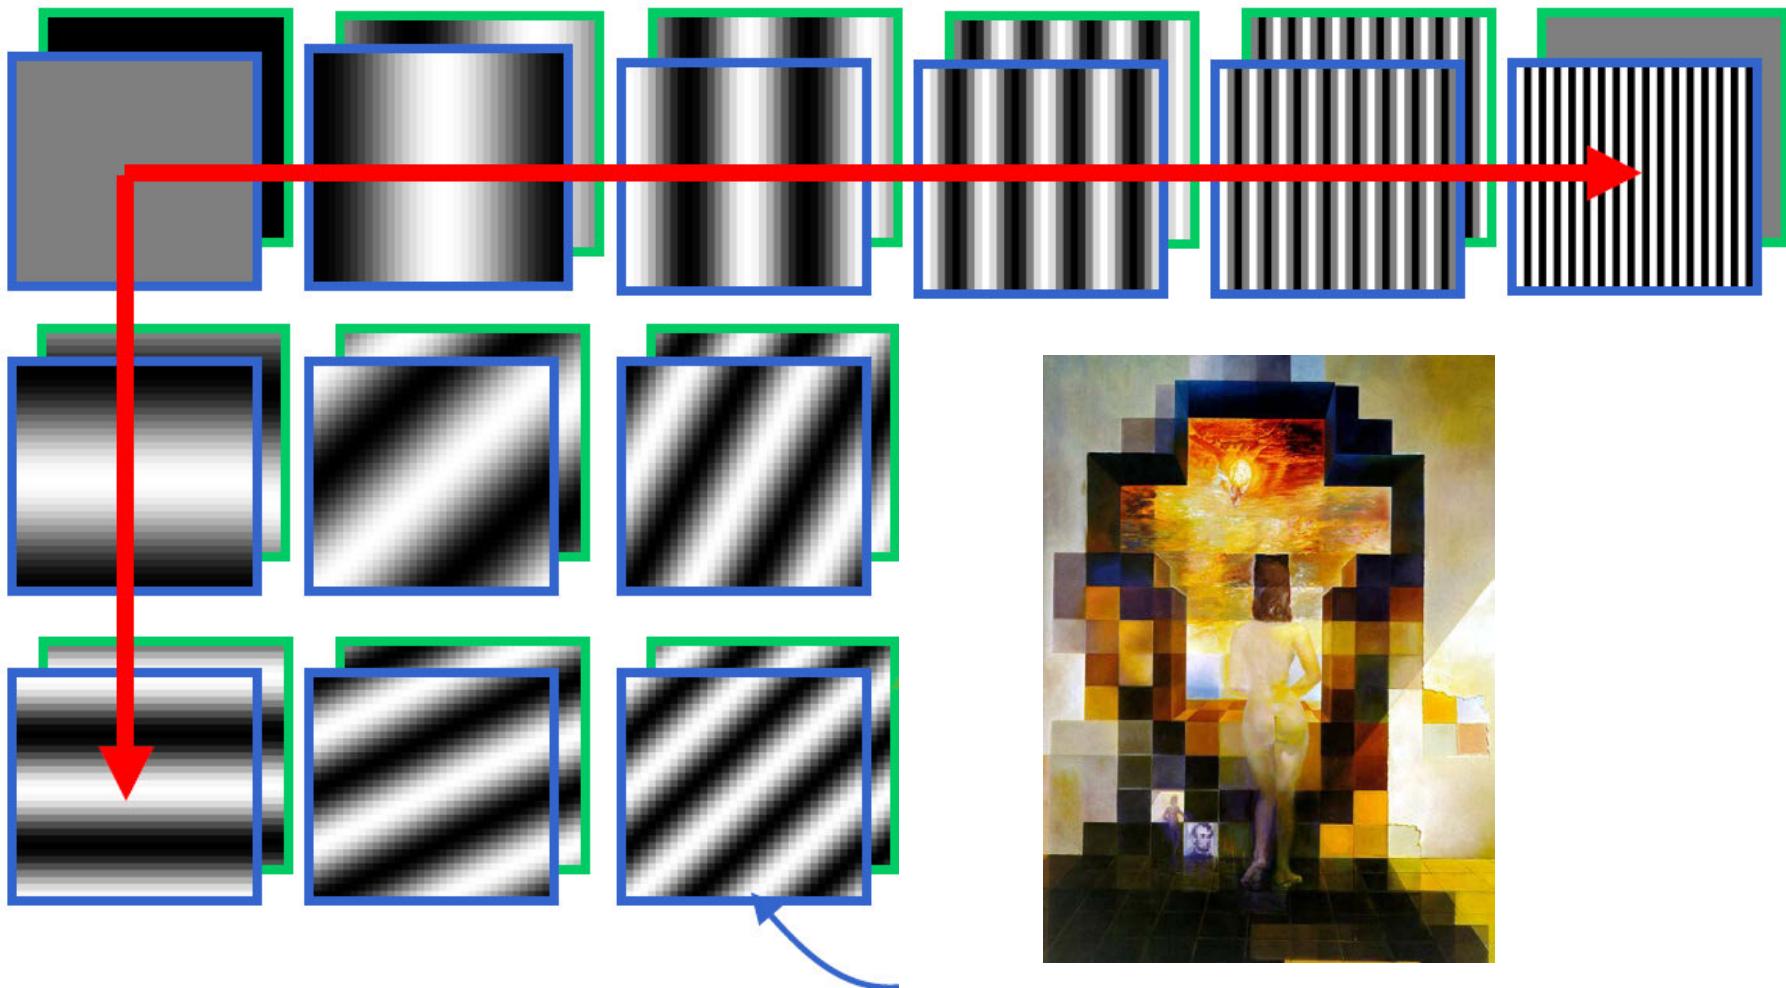






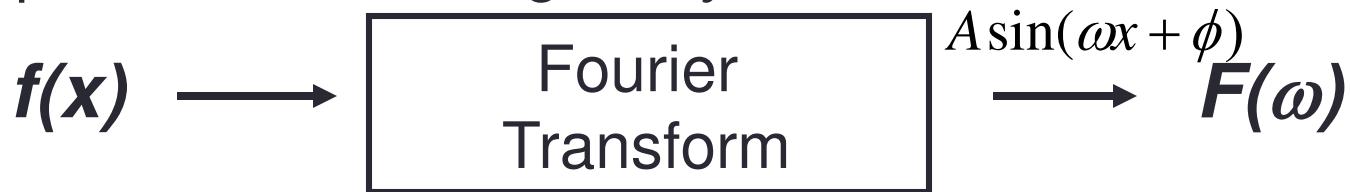
# Fourier: A nice set of basis

Teases away fast vs. slow changes in the image.



# Fourier Transform

We want to understand the frequency  $\omega$  of our signal. So, let's reparametrize the signal by  $\omega$  instead of  $x$ :



For every  $\omega$  from 0 to inf, (actually -inf to inf),  $F(\omega)$  holds the amplitude  $A$  and phase  $\phi$  of the corresponding sine

- How can  $F$  hold both? Complex number trick!

$$A = \pm \sqrt{R(\omega)^2 + I(\omega)^2}$$

$$F(\omega) = R(\omega) + iI(\omega)$$

*Even*      *Odd*

$$\phi = \tan^{-1} \frac{I(\omega)}{R(\omega)}$$

And we can go back:

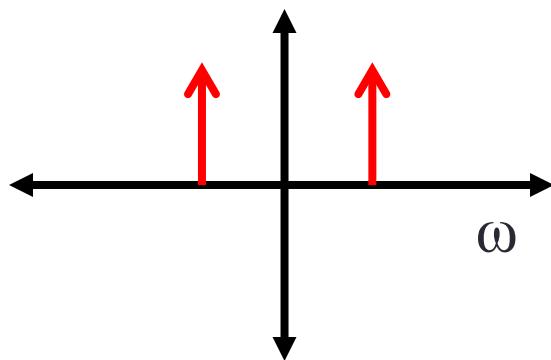


# Frequency Spectra – Even/Odd

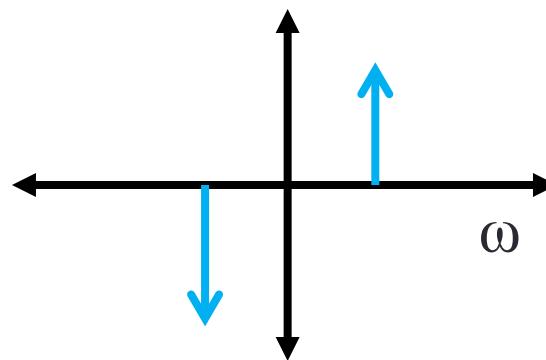
Frequency actually goes from  $-\infty$  to  $\infty$ .

Sinusoid example:

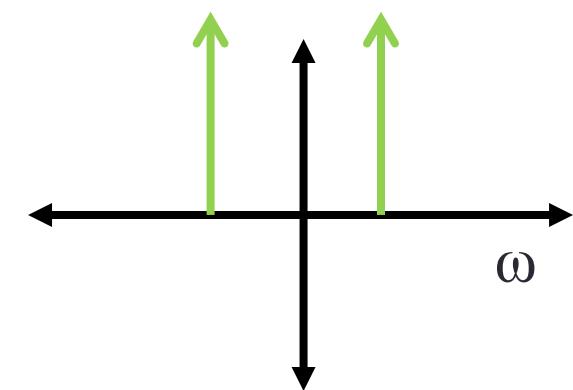
*Even (cos)*



*Odd (sin)*



*Magnitude*



*Real*

*Imaginary*

*Power*

# 2D Fourier Transforms

- The two dimensional version: .

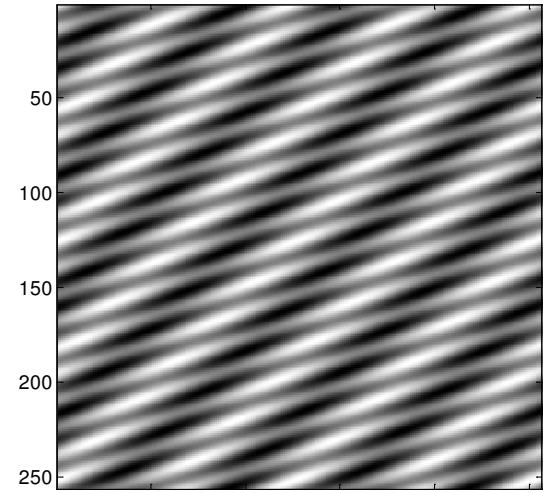
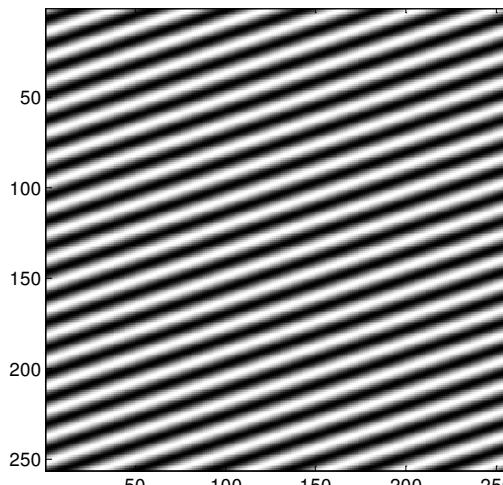
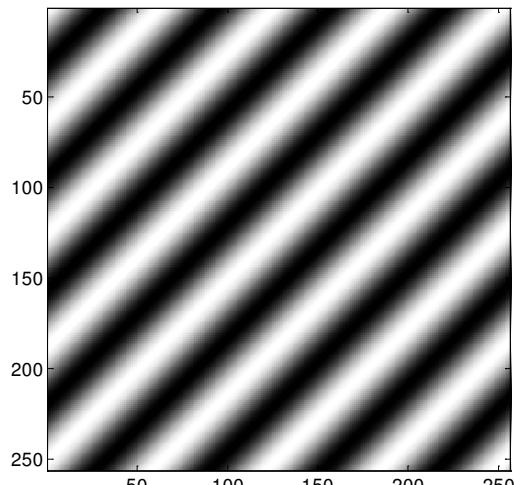
$$F(u, v) = \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(x, y) e^{-i 2\pi(ux+vy)} dx dy$$

- And the 2D ***Discrete FT***:

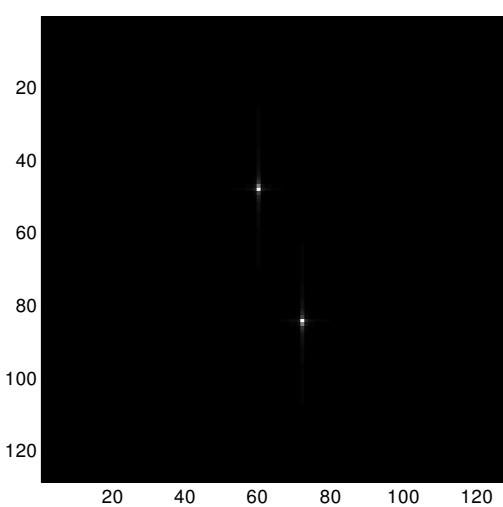
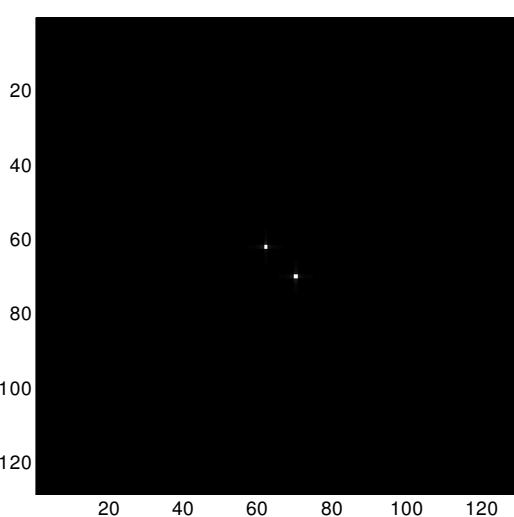
$$F(k_x, k_y) = \frac{1}{N} \sum_{x=0}^{N-1} \sum_{y=0}^{N-1} f(x, y) e^{-i \frac{2\pi(k_x x + k_y y)}{N}}$$

- Works best when you put the origin of  $k$  in the middle....

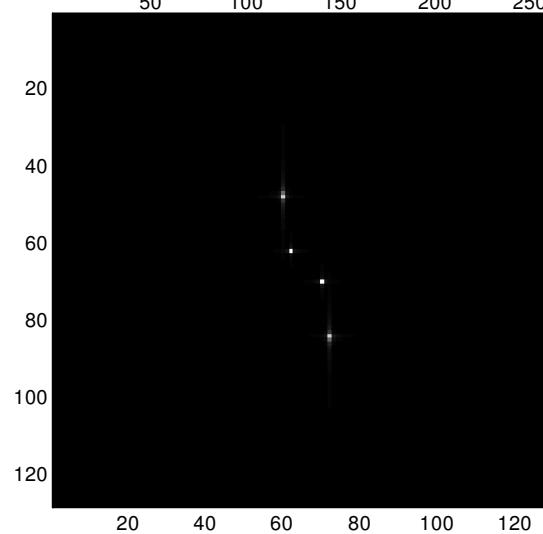
# Linearity of Sum



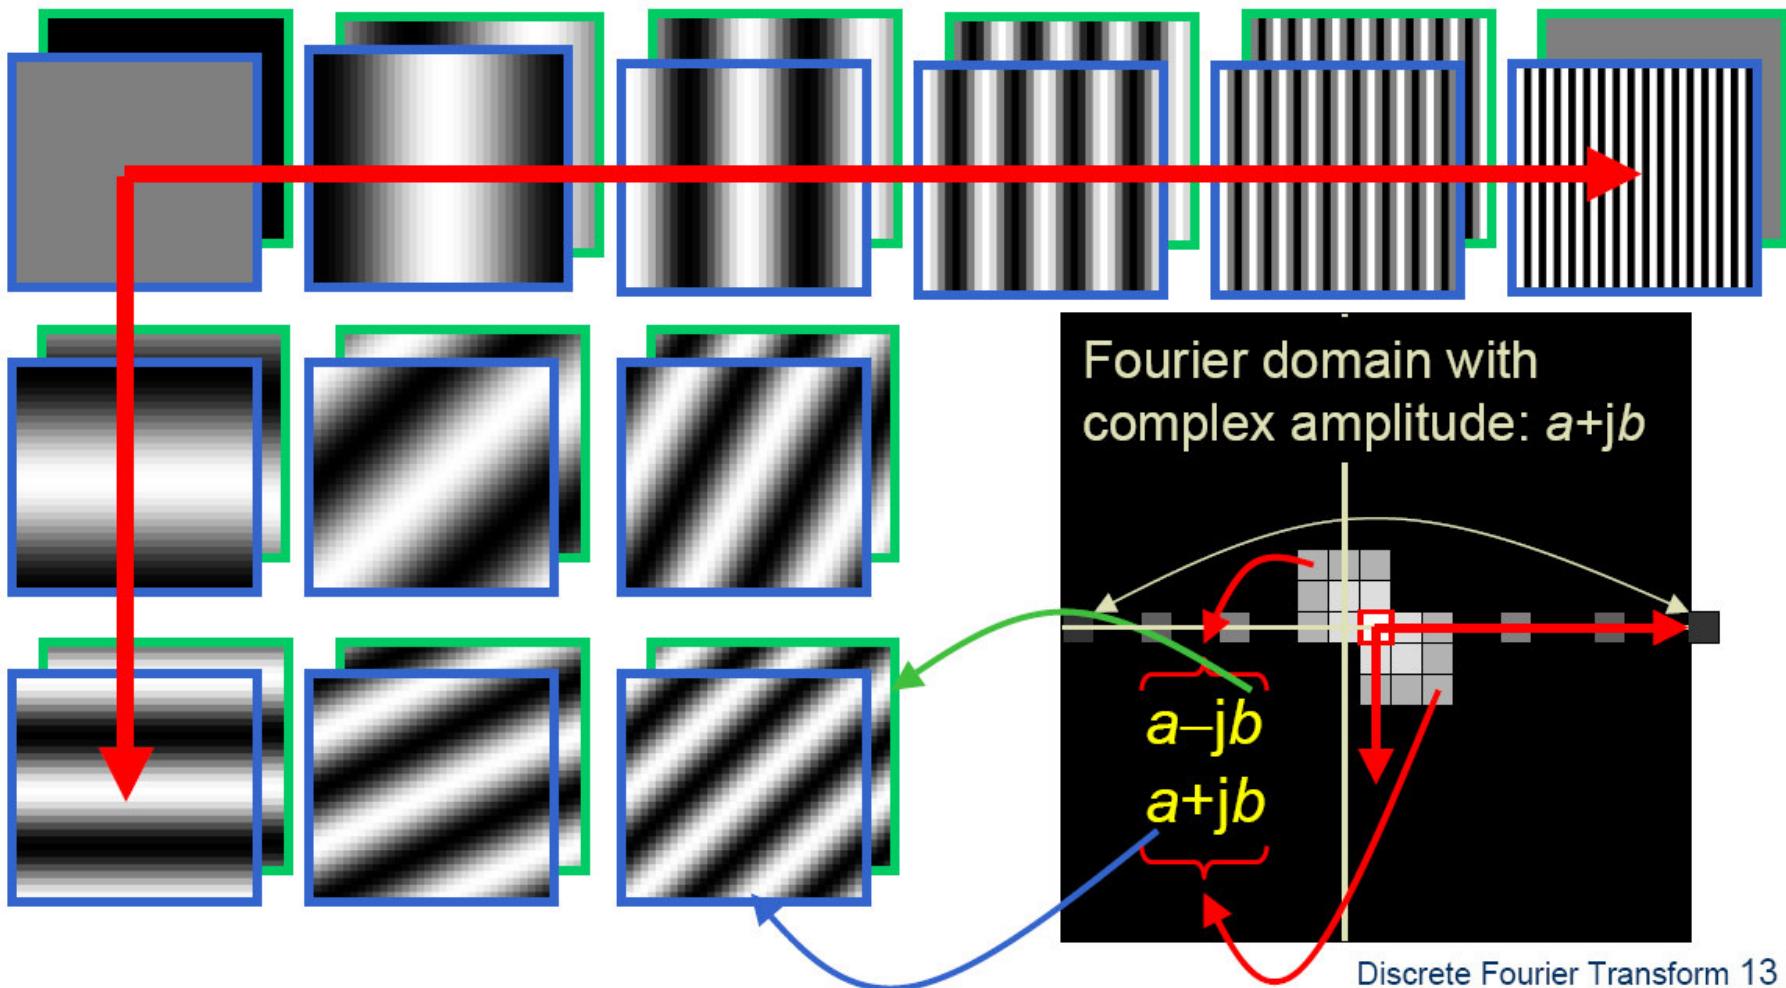
+



=

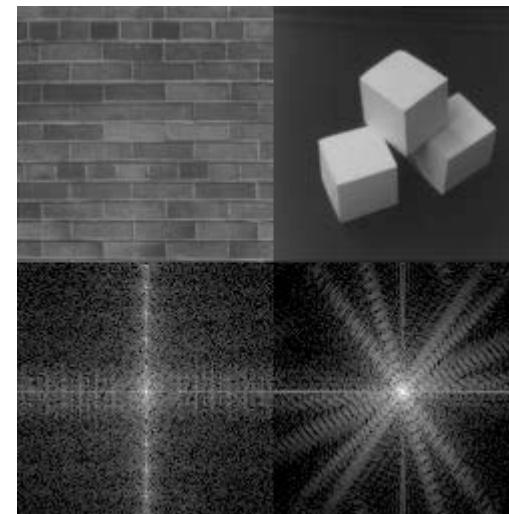
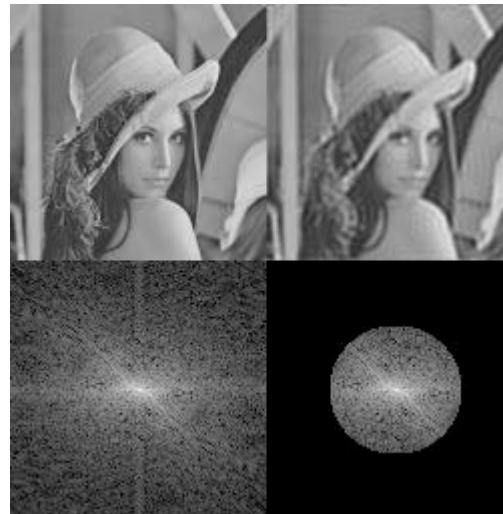
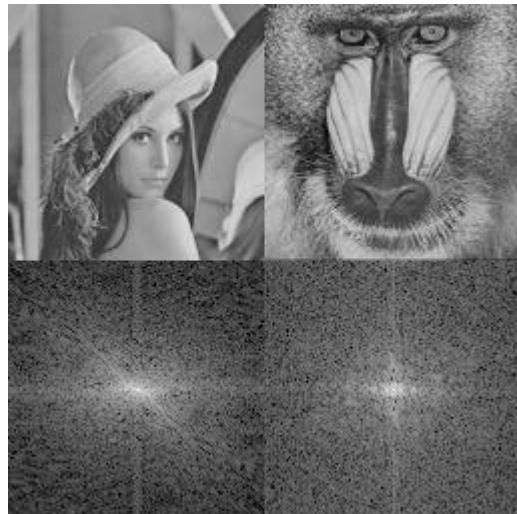


# Extension to 2D – Complex plane



Both a Real and Im version

# Examples



# Fourier Transform and Convolution

Let  $g = f * h$

Then  $G(u) = \int_{-\infty}^{\infty} g(x) e^{-i2\pi ux} dx$

$$= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} f(\tau) h(x - \tau) e^{-i2\pi ux} d\tau dx$$

$$= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} [f(\tau) e^{-i2\pi u\tau} d\tau] [h(x - \tau) e^{-i2\pi u(x - \tau)} dx]$$

$$= \int_{-\infty}^{\infty} [f(\tau) e^{-i2\pi u\tau} d\tau] \int_{-\infty}^{\infty} [h(x') e^{-i2\pi u x'} dx']$$

$$= F(u) H(u)$$

*Convolution in spatial domain*

$\Leftrightarrow$  *Multiplication in frequency domain*

# Fourier Transform and Convolution

Spatial Domain ( $x$ )	$\leftrightarrow$	Frequency Domain ( $u$ )
$g = f * h$	$\leftrightarrow$	$G = FH$
$g = fh$	$\leftrightarrow$	$G = F * H$

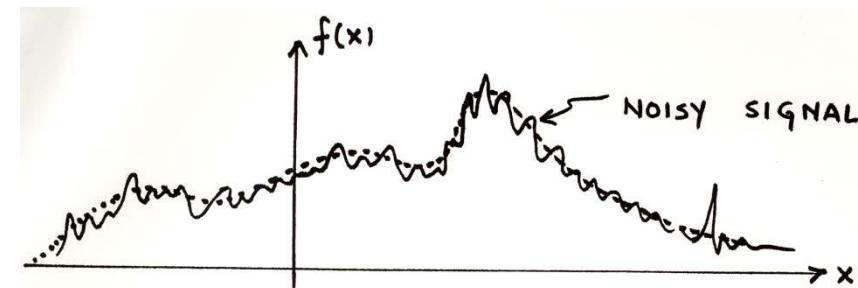
So, we can find  $g(x)$  by Fourier transform

$$\begin{array}{c} g \\ \uparrow \\ \boxed{\text{IFT}} \\ G \end{array} = \begin{array}{c} f \\ \downarrow \\ \boxed{\text{FT}} \\ F \end{array} * \begin{array}{c} h \\ \downarrow \\ \boxed{\text{FT}} \\ H \end{array}$$

# Example use: Smoothing/Blurring

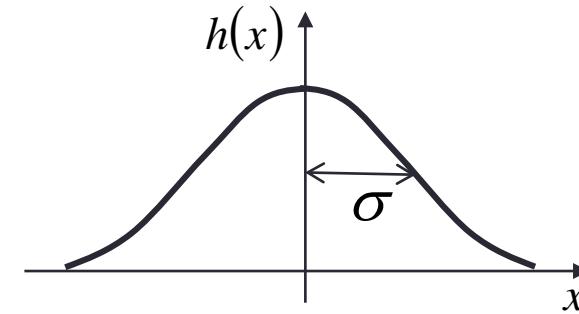
- We want a smoothed function of  $f(x)$

$$g(x) = f(x) * h(x)$$



- Let us use a Gaussian kernel

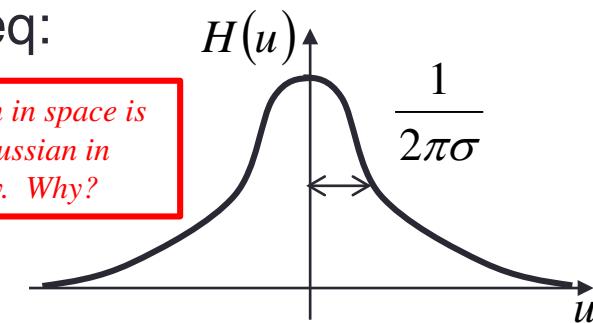
$$h(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left[-\frac{1}{2} \frac{x^2}{\sigma^2}\right]$$



- Convolution in space is multiplication in freq:

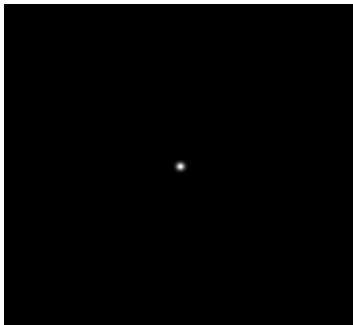
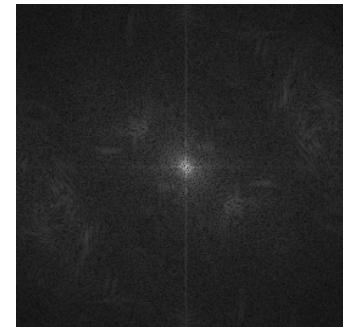
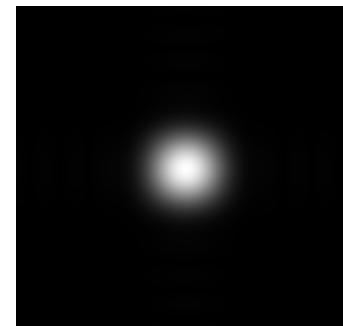
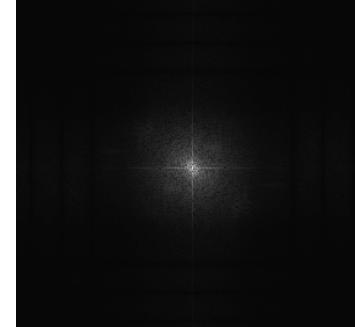
$$G(u) = F(u)H(u)$$

*Fat Gaussian in space is skinny Gaussian in frequency. Why?*

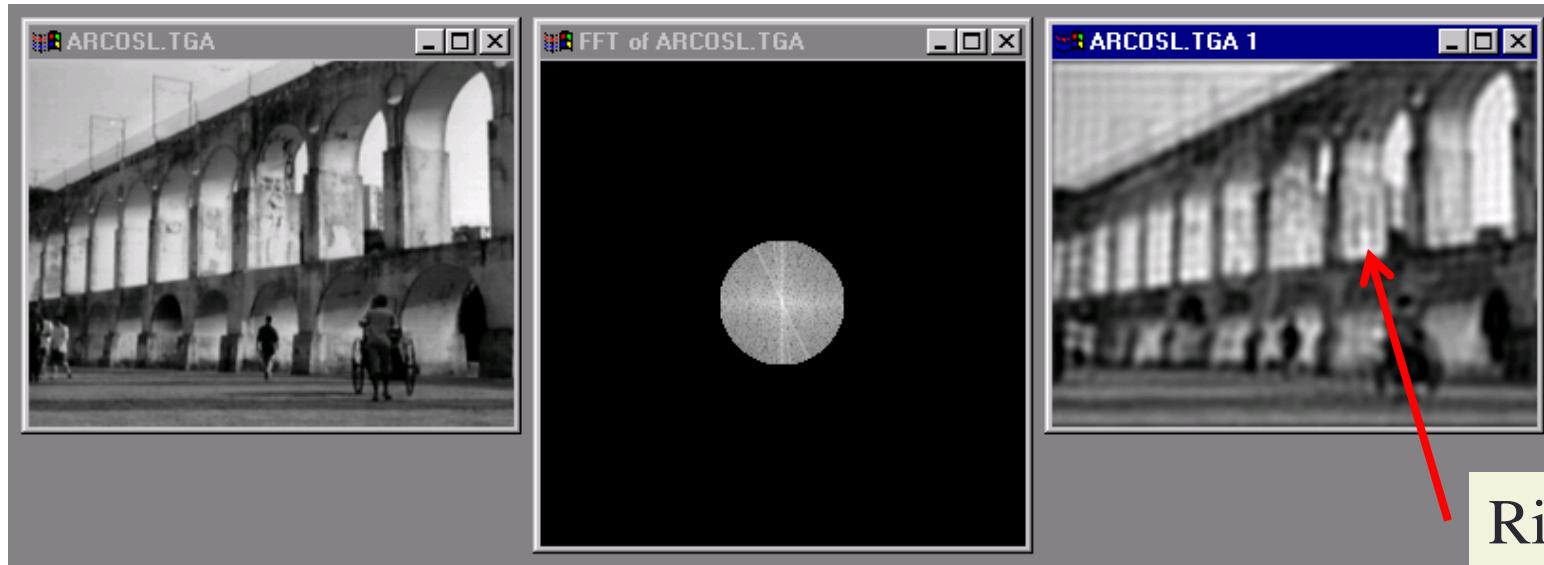


$H(u)$  **attenuates** high frequencies in  $F(u)$  (Low-pass Filter)!

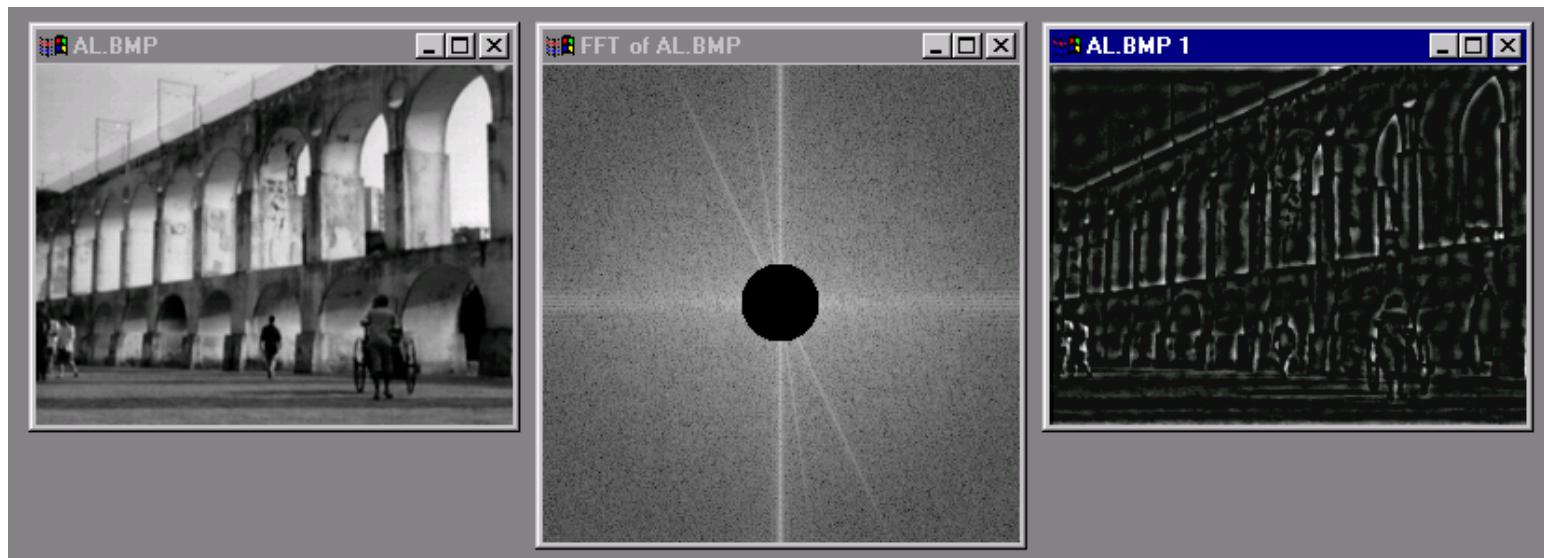
# 2D convolution theorem example

 $f(x,y)$  $*$  $h(x,y)$  $\Downarrow$  $g(x,y)$  $\times$  $|F(s_x, s_y)|$  $(\text{or } |F(u, v)|)$  $\Downarrow$  $|H(s_x, s_y)|$  $|G(s_x, s_y)|$

# Low and High Pass filtering

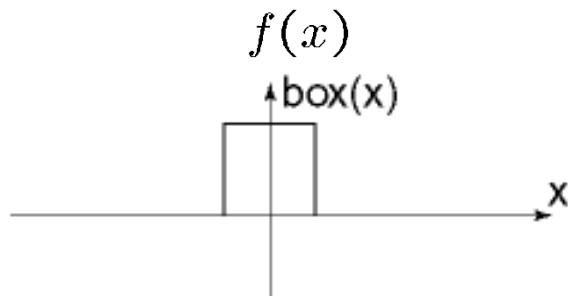


Ringing

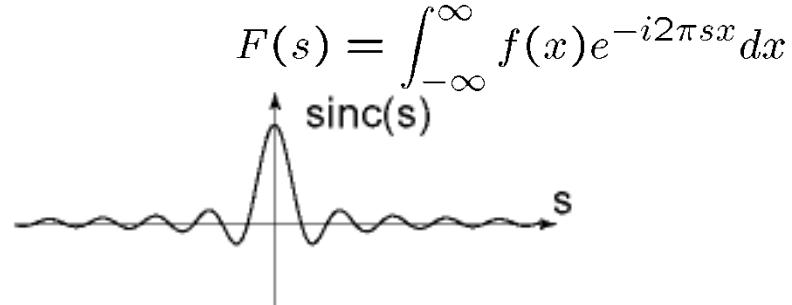


# Fourier Transform smoothing pairs

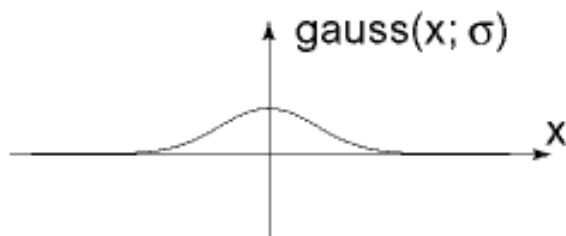
Spatial domain



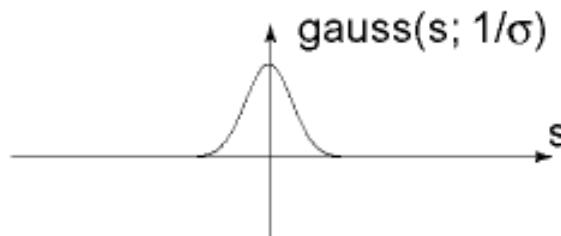
Frequency domain



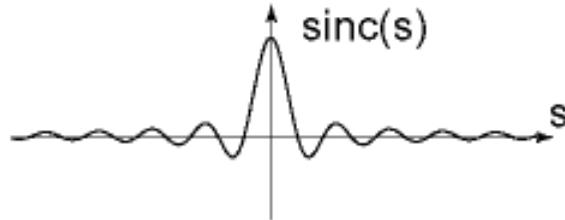
$\text{gauss}(x; \sigma)$



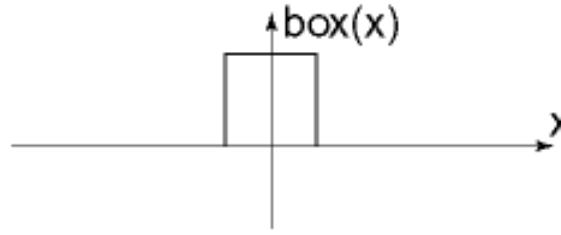
$\text{gauss}(s; 1/\sigma)$



$\text{sinc}(s)$



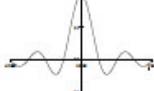
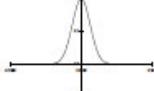
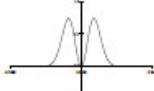
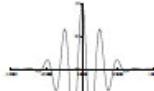
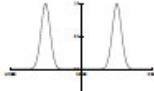
$\text{box}(x)$



# Properties of Fourier Transform

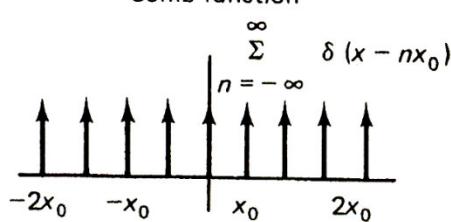
	Spatial Domain ( $x$ )	Frequency Domain ( $u$ )
<b>Linearity</b>	$c_1 f(x) + c_2 g(x)$	$c_1 F(u) + c_2 G(u)$
<b>Scaling</b>	$f(ax)$	$\frac{1}{ a } F\left(\frac{u}{a}\right)$
<b>Shifting</b>	$f(x - x_0)$	$e^{-i2\pi u x_0} F(u)$
<b>Symmetry</b>	$F(x)$	$f(-u)$
<b>Conjugation</b>	$f^*(x)$	$F^*(-u)$
<b>Convolution</b>	$f(x) * g(x)$	$F(u)G(u)$
<b>Differentiation</b>	$\frac{d^n f(x)}{dx^n}$	$(i2\pi u)^n F(u)$
		<b>Multiply by <math>u</math></b>

# Fourier Pairs (from Szeliski)

Name	Signal		Transform
impulse		$\delta(x) \Leftrightarrow 1$	
shifted impulse		$\delta(x - u) \Leftrightarrow e^{-j\omega u}$	
box filter		$\text{box}(x/a) \Leftrightarrow \text{asinc}(a\omega)$	
tent		$\text{tent}(x/a) \Leftrightarrow \text{asinc}^2(a\omega)$	
Gaussian		$G(x; \sigma) \Leftrightarrow \frac{\sqrt{2\pi}}{\sigma} G(\omega; \sigma^{-1})$	
Laplacian of Gaussian		$(\frac{x^2}{\sigma^4} - \frac{1}{\sigma^2})G(x; \sigma) \Leftrightarrow -\frac{\sqrt{2\pi}}{\sigma}\omega^2 G(\omega; \sigma^{-1})$	
Gabor		$\cos(\omega_0 x)G(x; \sigma) \Leftrightarrow \frac{\sqrt{2\pi}}{\sigma} G(\omega \pm \omega_0; \sigma^{-1})$	
unsharp mask		$(1 + \gamma)\delta(x) - \gamma G(x; \sigma) \Leftrightarrow \frac{(1 + \gamma)}{\sigma} - \frac{\sqrt{2\pi}\gamma}{\sigma} G(\omega; \sigma^{-1})$	
windowed sinc		$r\cos(x/(aW)) \quad \text{sinc}(x/a) \Leftrightarrow \text{(see Figure 3.29)}$	

# Fourier Transform Sampling Pairs

Comb function



FT of an “impulse train”  
is an impulse train

$$\frac{1}{x_0} \sum_{n=-\infty}^{\infty} \delta(\xi - \frac{n}{x_0})$$

A diagram showing a series of vertical arrows on a horizontal axis. The arrows are positioned at regular intervals along the axis. Above the diagram, the equation  $\frac{1}{x_0} \sum_{n=-\infty}^{\infty} \delta(\xi - \frac{n}{x_0})$  is written. Below the axis, the labels  $-\frac{2}{x_0}$ ,  $-\frac{1}{x_0}$ ,  $\frac{1}{x_0}$ , and  $\frac{2}{x_0}$  are placed under the corresponding arrows.

$$\cos 2\pi\omega_0 x$$

A graph of a periodic wave oscillating above and below a horizontal axis. The wave has a single peak above the axis and one trough below it. The label  $\cos 2\pi\omega_0 x$  is written above the peak.

$$\frac{1}{2} [\delta(\xi - \omega_0) + \delta(\xi + \omega_0)]$$

A graph of a periodic wave oscillating above and below a horizontal axis. It consists of two vertical arrows, one pointing up at  $\omega_0$  and one pointing down at  $-\omega_0$ . The label  $\frac{1}{2} [\delta(\xi - \omega_0) + \delta(\xi + \omega_0)]$  is written above the positive arrow.

$$\sin 2\pi\omega_0 x$$

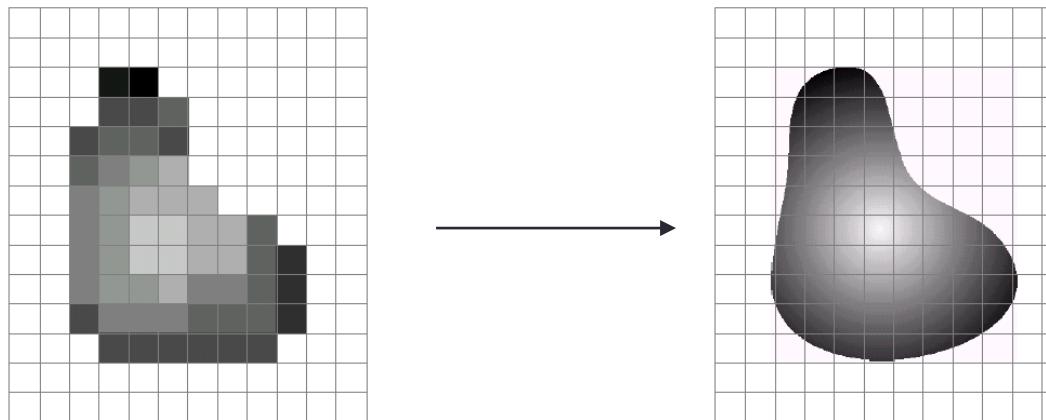
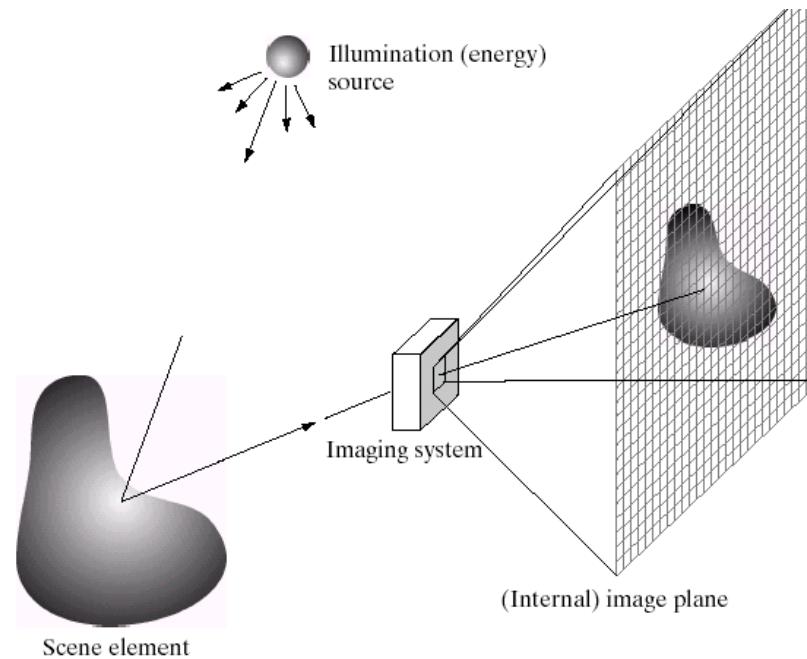
A graph of a periodic wave oscillating above and below a horizontal axis. It has a single zero-crossing point on the axis. The label  $\sin 2\pi\omega_0 x$  is written above the zero-crossing point.

$$\frac{1}{2} i [-\delta(\xi - \omega_0) + \delta(\xi + \omega_0)]$$

A graph of a periodic wave oscillating above and below a horizontal axis. It consists of two vertical arrows, one pointing up at  $\omega_0$  and one pointing down at  $-\omega_0$ . The label  $\frac{1}{2} i [-\delta(\xi - \omega_0) + \delta(\xi + \omega_0)]$  is written above the positive arrow. To the right of the graph, the label "Im F" is written, indicating the imaginary part of the Fourier transform.

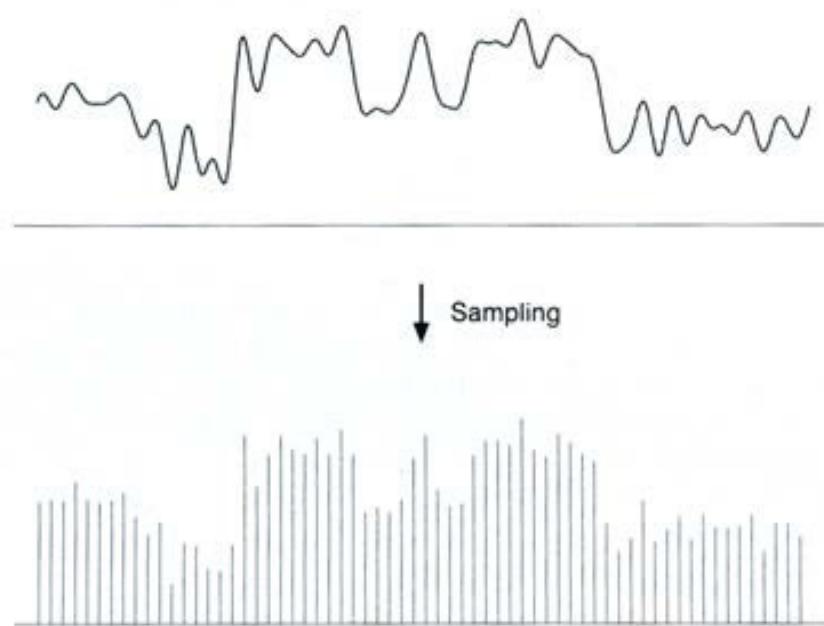
# Sampling and Aliasing

# Sampling and Reconstruction



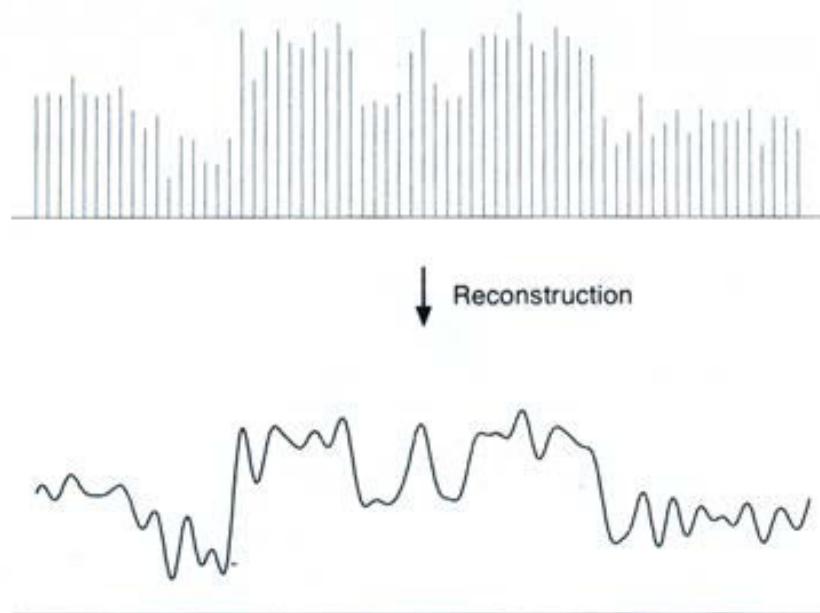
# Sampled representations

- How to store and compute with continuous functions?
- Common scheme for representation: samples
  - write down the function's values at many points

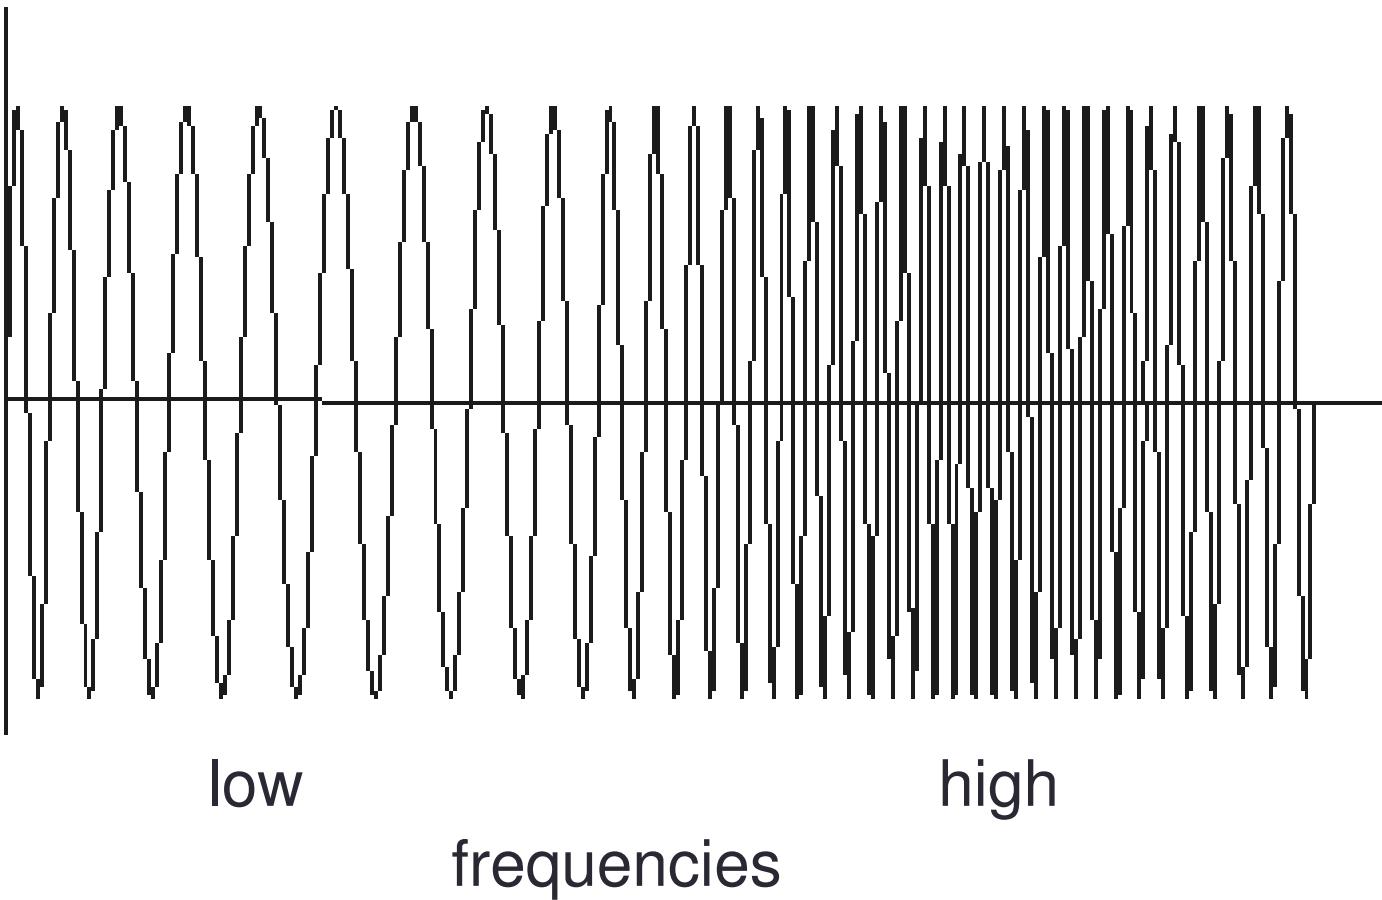


# Reconstruction

- Making samples back into a continuous function
  - for output (need realizable method)
  - for analysis or processing (need mathematical method)
  - amounts to “guessing” what the function did in between

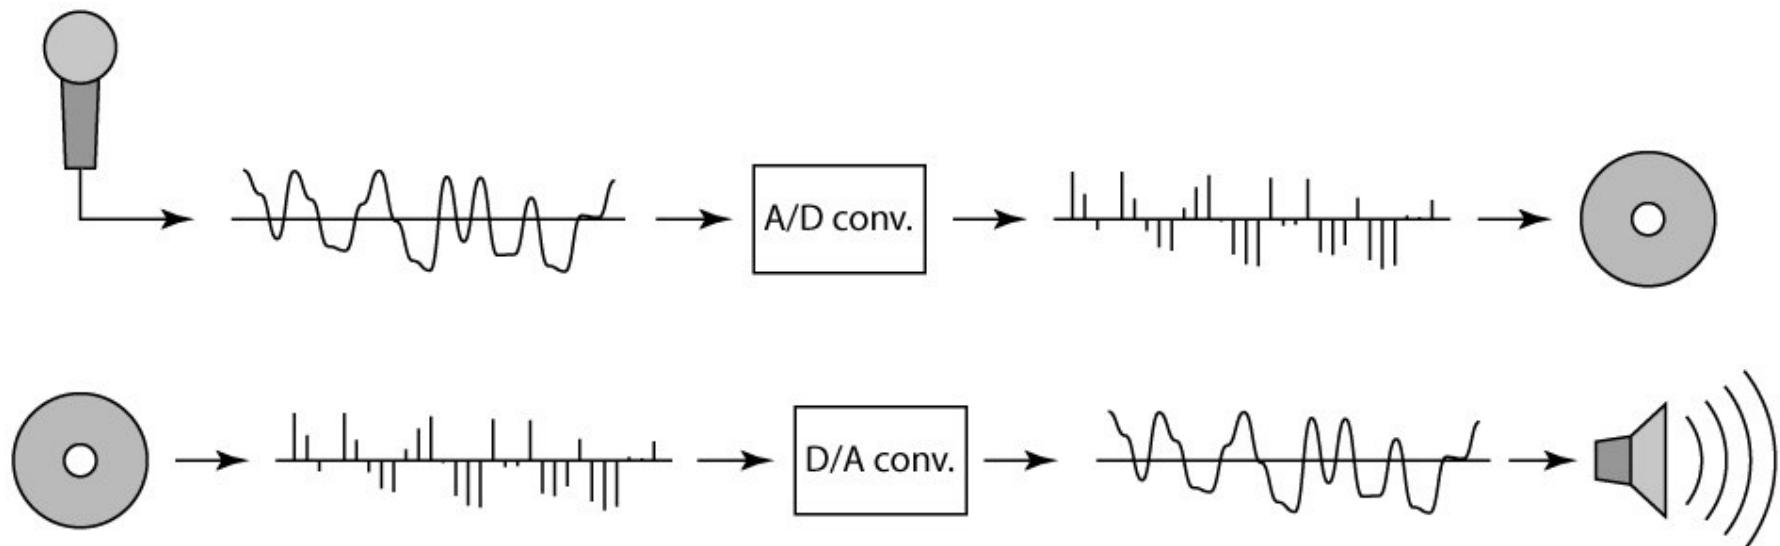


# 1D Example: Audio



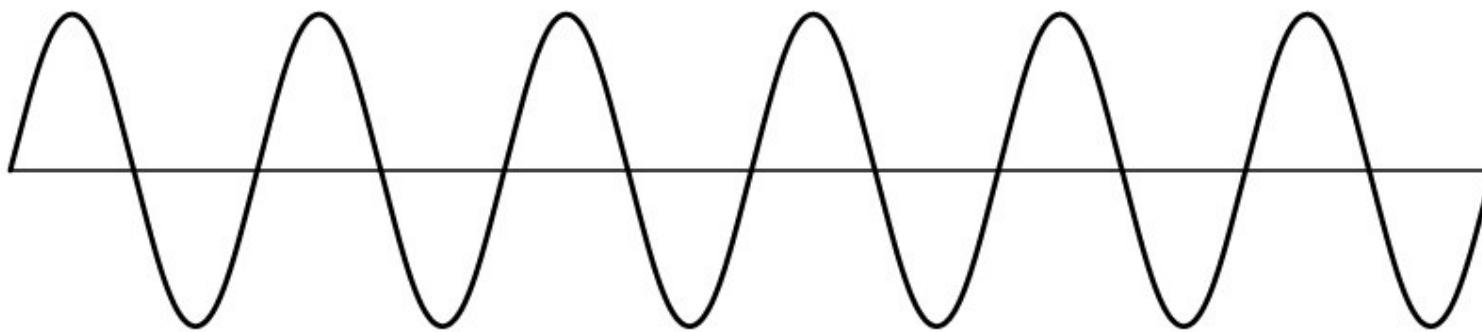
# Sampling in digital audio

- Recording: sound to analog to samples to disc
- Playback: disc to samples to analog to sound again
  - how can we be sure we are filling in the gaps correctly?



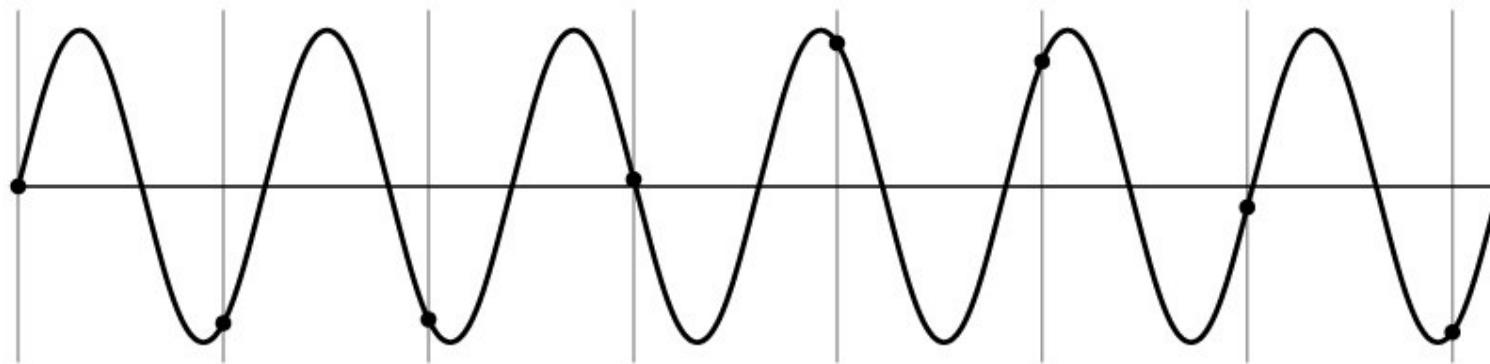
# Sampling and Reconstruction

- Simple example: a sign wave



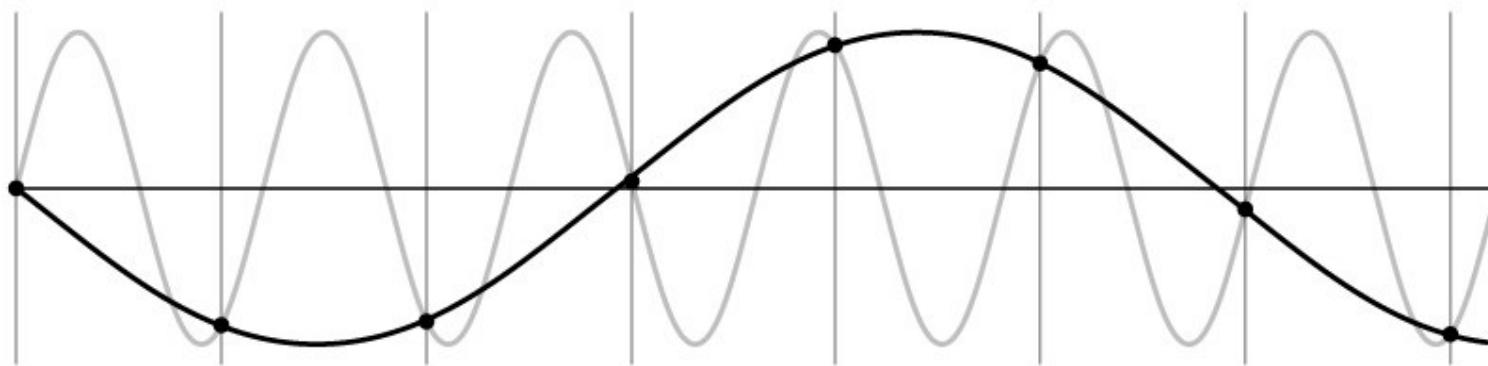
# Undersampling

- What if we “missed” things between the samples?
- Simple example: undersampling a sine wave
  - unsurprising result: information is lost



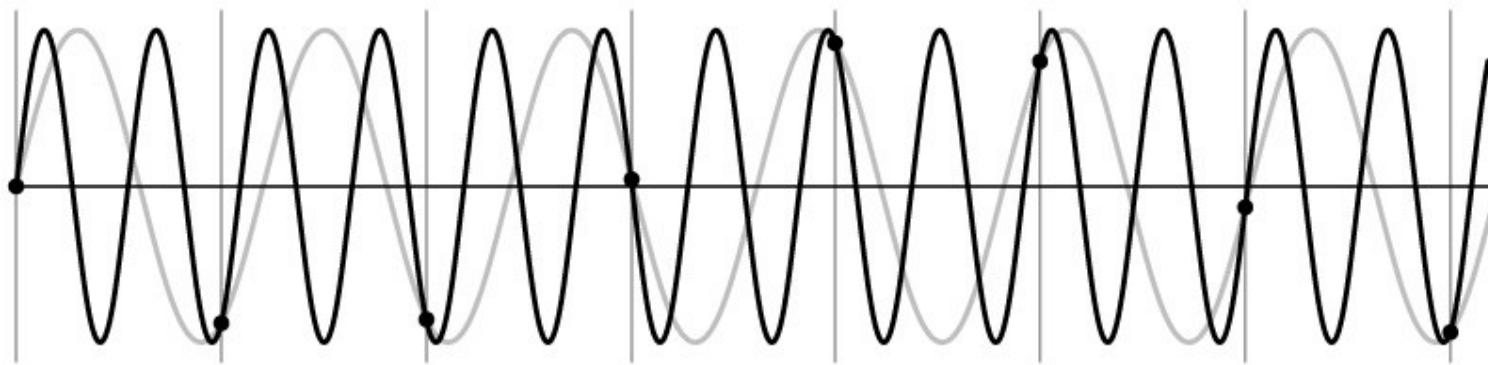
# Undersampling

- What if we “missed” things between the samples?
- Simple example: undersampling a sine wave
  - unsurprising result: information is lost
  - surprising result: indistinguishable from lower frequency



# Undersampling

- What if we “missed” things between the samples?
- Simple example: undersampling a sine wave
  - unsurprising result: information is lost
  - surprising result: indistinguishable from lower frequency
  - also was always indistinguishable from higher frequencies
  - aliasing: signals “traveling in disguise” as other frequencies

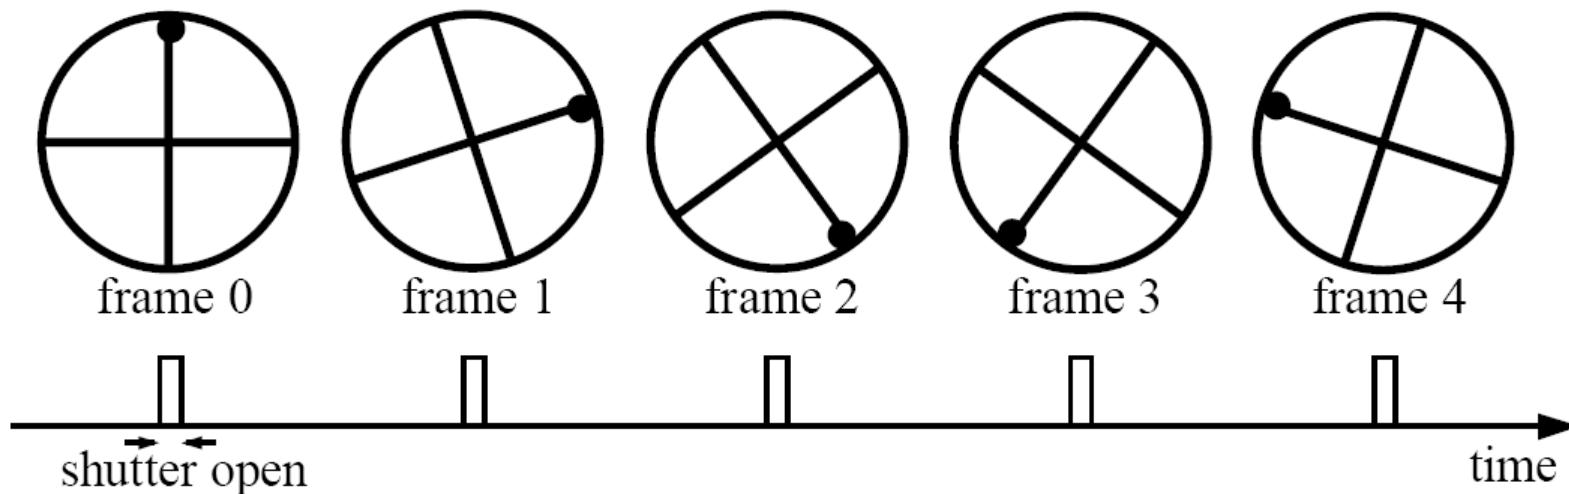


# Aliasing in video

Imagine a spoked wheel moving to the right (rotating clockwise).

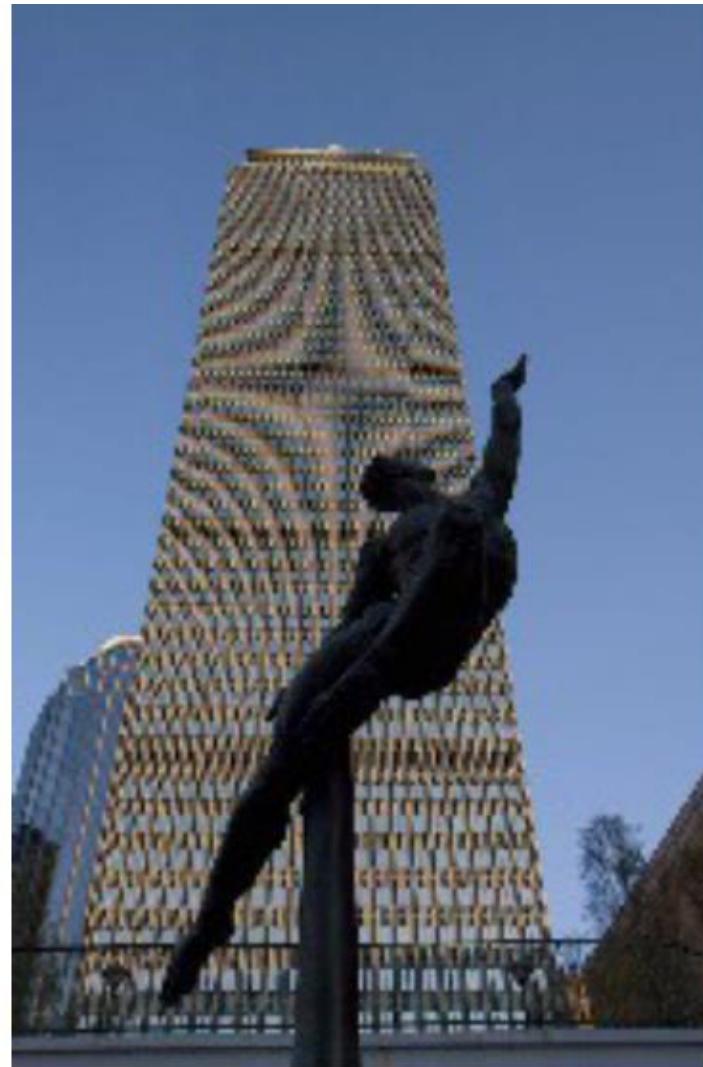
Mark wheel with dot so we can see what's happening.

If camera shutter is only open for a fraction of a frame time (frame time = 1/30 sec. for video, 1/24 sec. for film):



Without dot, wheel appears to be rotating slowly backwards!  
(counterclockwise)

# Image sub-sampling



Source: F. Durand

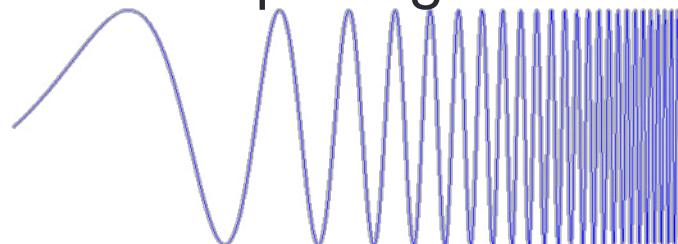
# Aliasing in images



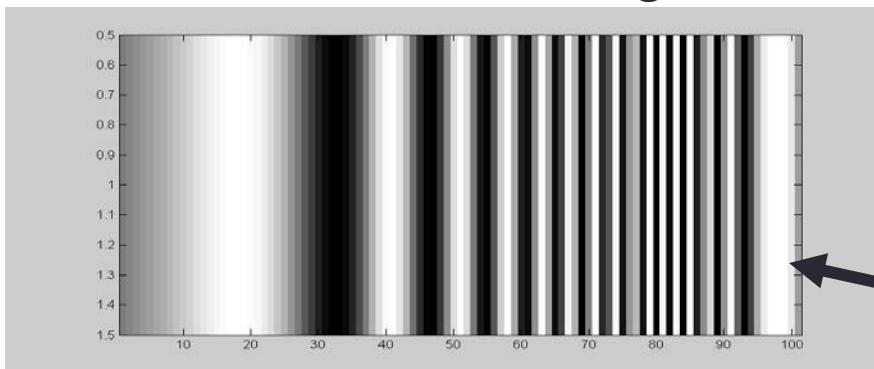
**Disintegrating textures**

# What's happening?

Input signal:

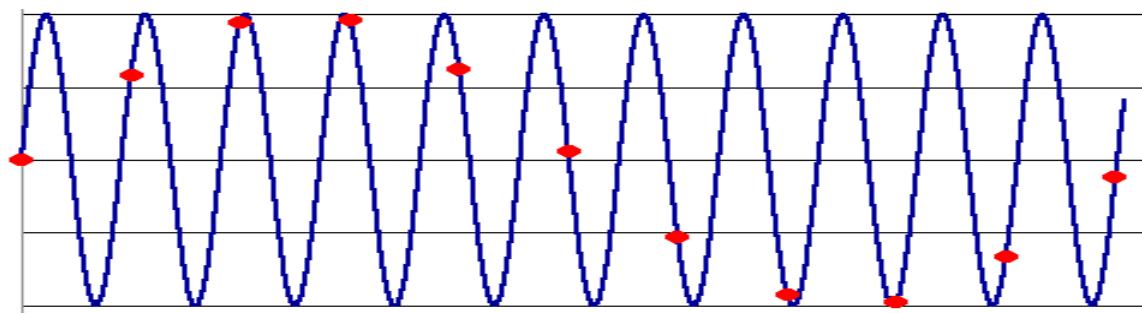


Plot as image:

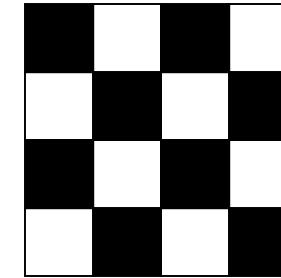
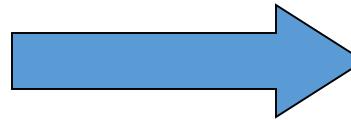
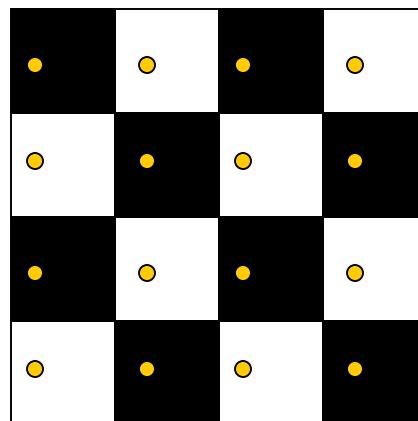
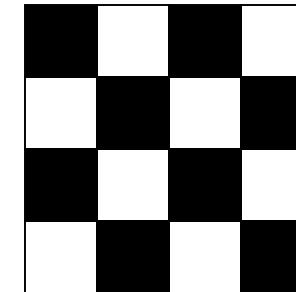
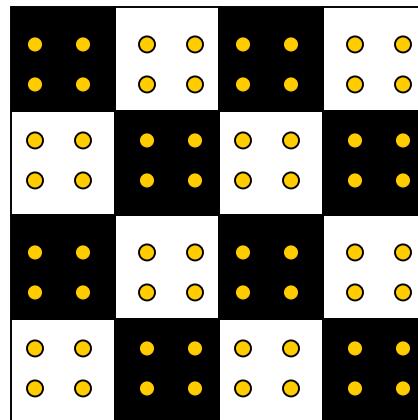


Alias!  
Not enough samples

```
x = 0:.05:5; imagesc(sin((2.^x).*x))
```

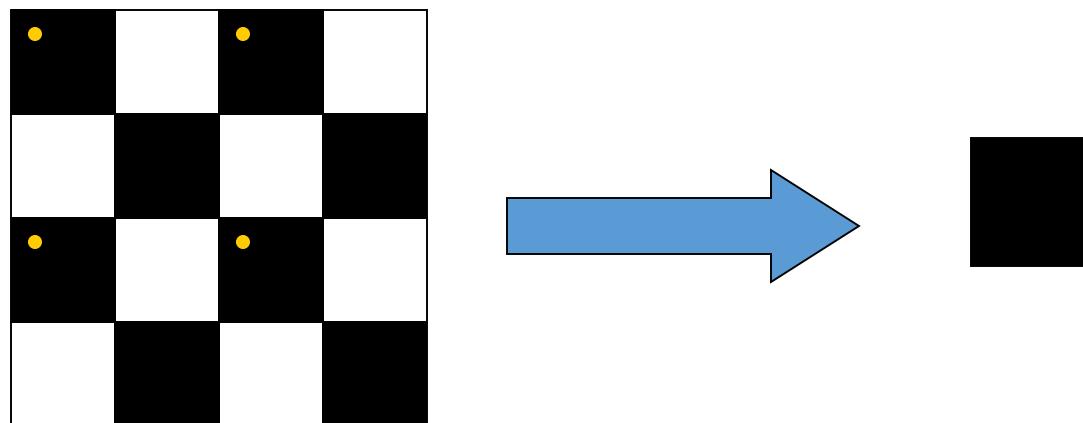
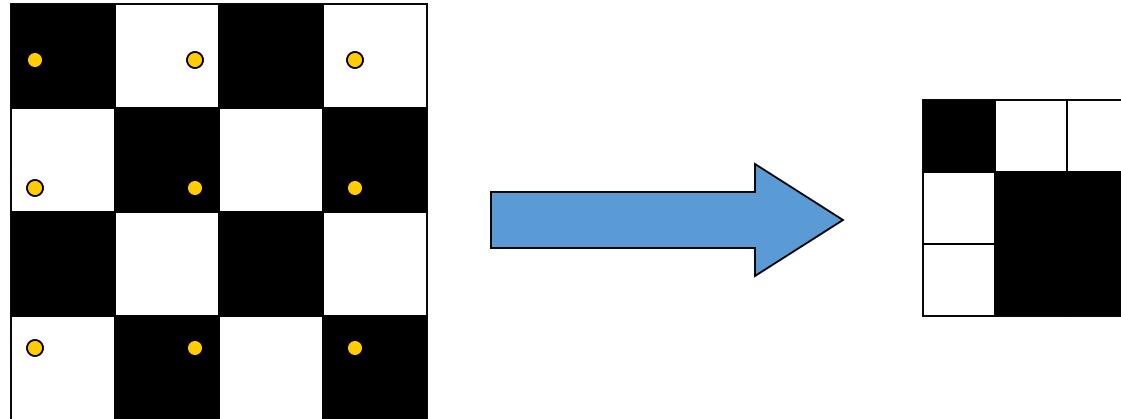


# Sampling an image



Examples of GOOD sampling

# Undersampling



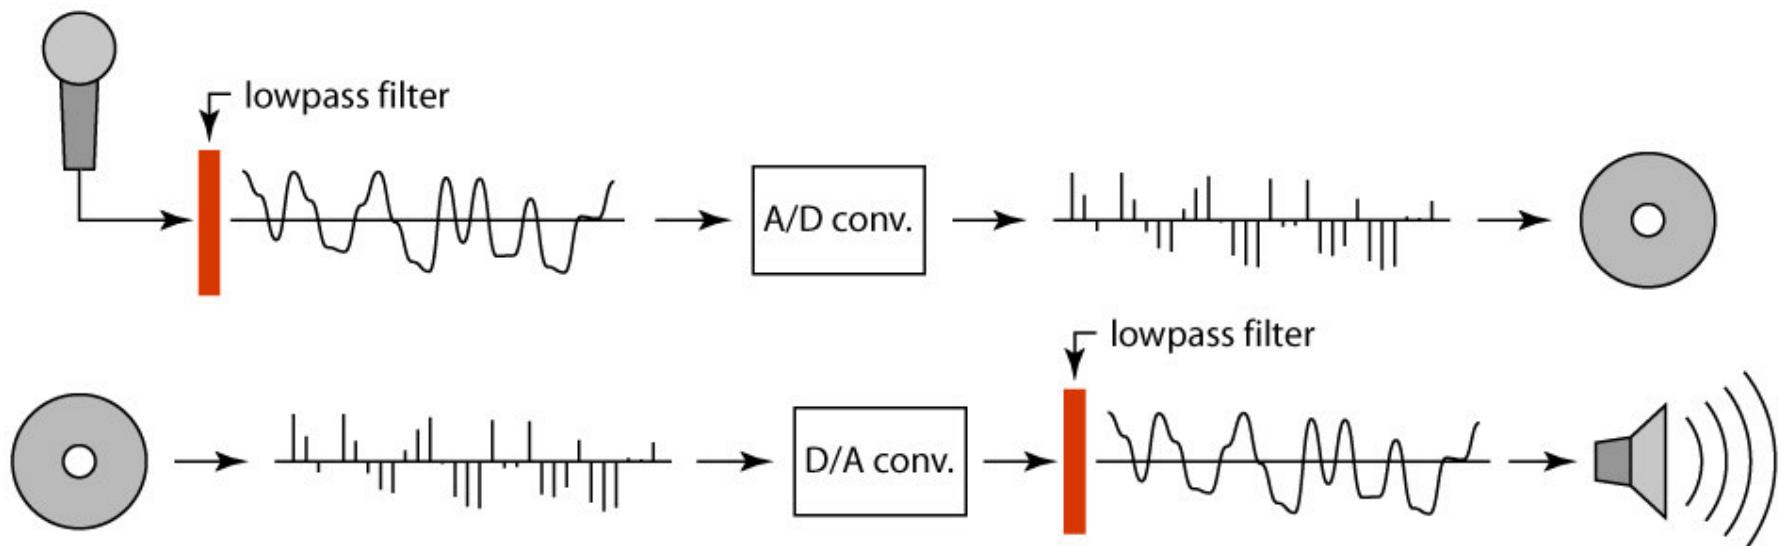
Examples of BAD sampling -> Aliasing

# Antialiasing

- What can we do about aliasing?
- Sample more often
  - Join the Mega-Pixel craze of the photo industry
  - But this can't go on forever
- Make the signal less “wiggly”
  - Get rid of some high frequencies
  - Will lose information
  - But it's better than aliasing

# Preventing aliasing

- Introduce lowpass filters:
  - remove high frequencies leaving only safe, low frequencies
  - choose lowest frequency in reconstruction (disambiguate)

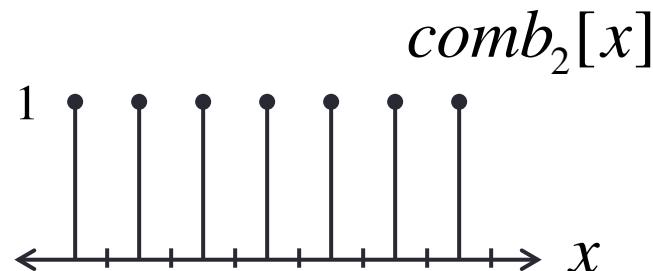


# Impulse Train

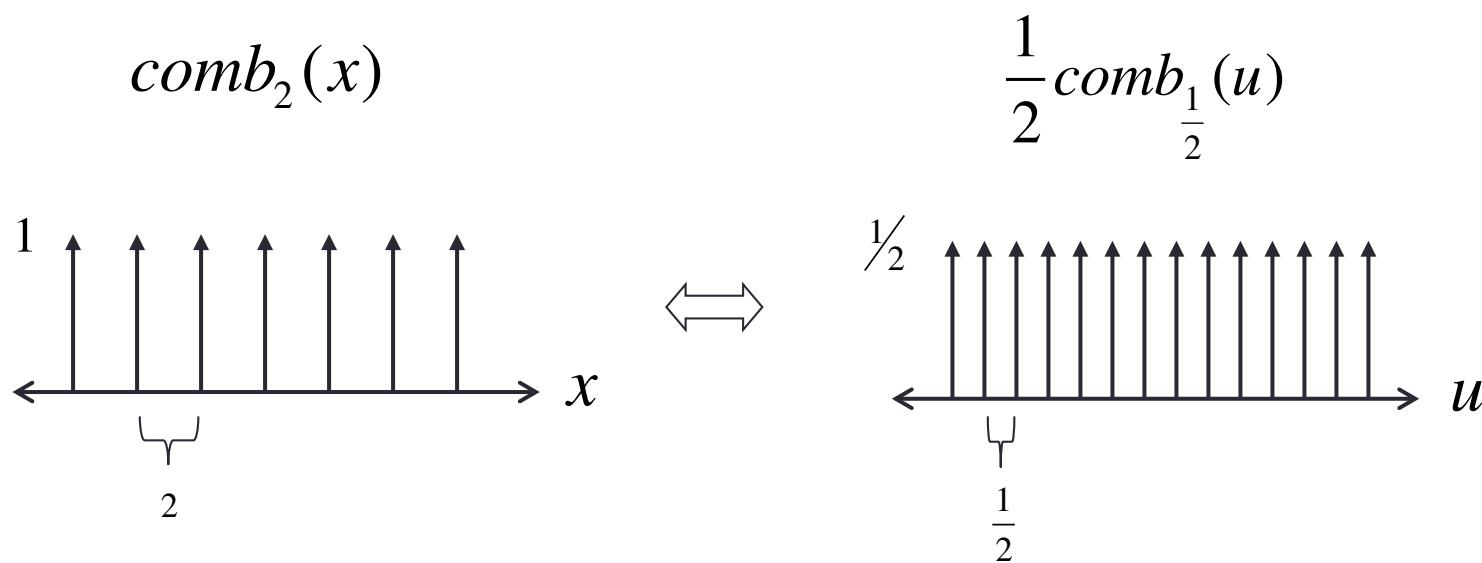
- Define a *comb* function (impulse train) in 1D as follows

$$\text{comb}_M[x] = \sum_{k=-\infty}^{\infty} \delta[x - kM]$$

where  $M$  is an integer



# Impulse Train in 1D



*Remember:*

**Scaling**

$$f(ax)$$

$$\frac{1}{|a|} F\left(\frac{u}{a}\right)$$

# Impulse Train in 2D (*bed of nails*)

$$\text{comb}_{M,N}(x, y) \equiv \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} \delta(x - kM, y - lN)$$

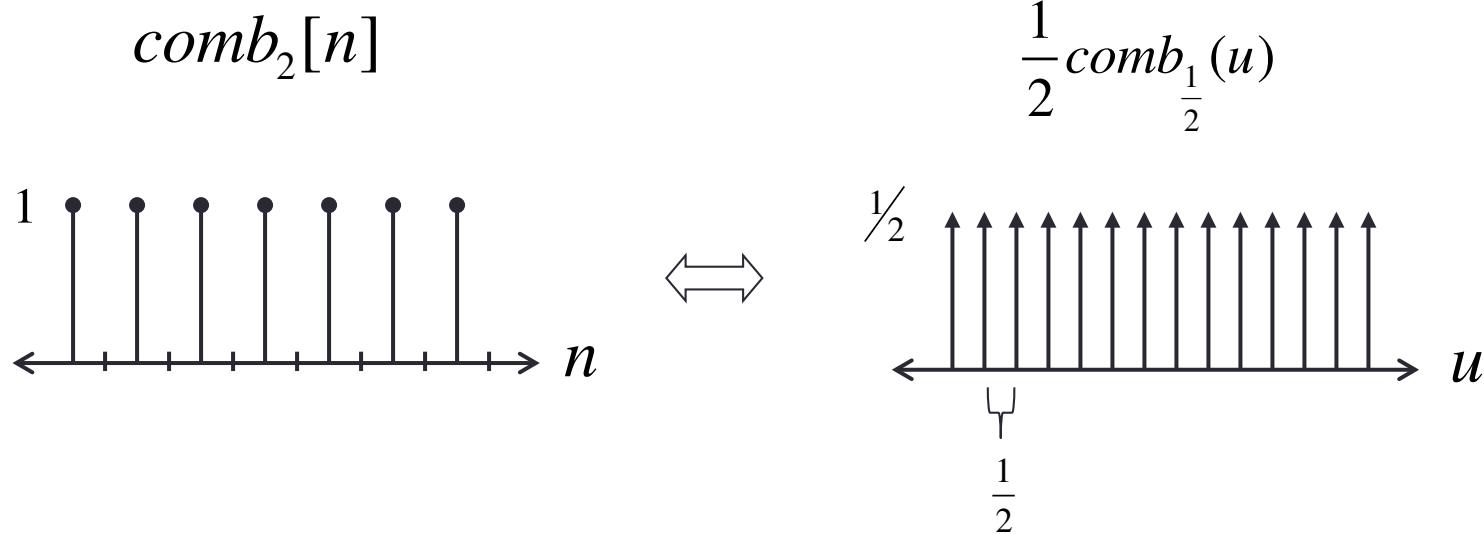
- Fourier Transform of an impulse train is also an impulse train:

$$\sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} \delta(x - kM, y - lN) \Leftrightarrow \frac{1}{MN} \sum_{k=-\infty}^{\infty} \sum_{l=-\infty}^{\infty} \delta\left(u - \frac{k}{M}, v - \frac{l}{N}\right)$$

$\text{comb}_{M,N}(x, y)$ 
 $\text{comb}_{\frac{1}{M}, \frac{1}{N}}(u, v)$

*As the comb samples get further apart, the spectrum samples get closer together!*

# Impulse Train



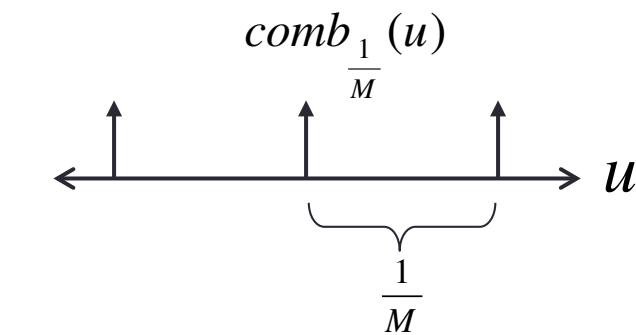
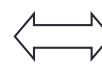
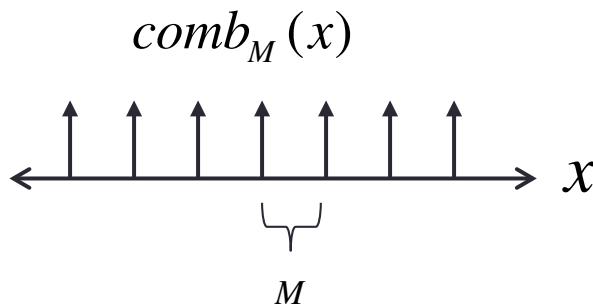
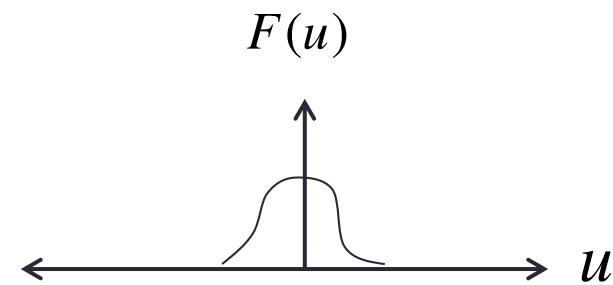
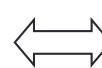
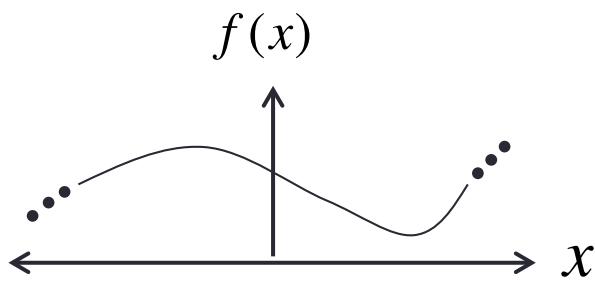
*Remember:*

**Scaling**

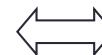
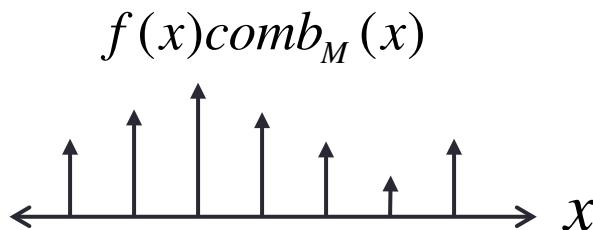
$$f(ax)$$

$$\frac{1}{|a|} F\left(\frac{u}{a}\right)$$

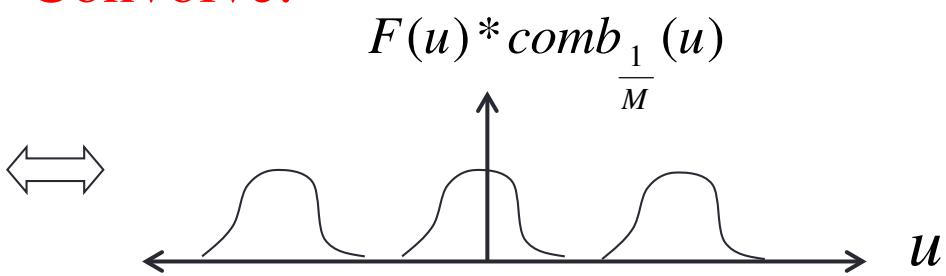
# Sampling low frequency signal



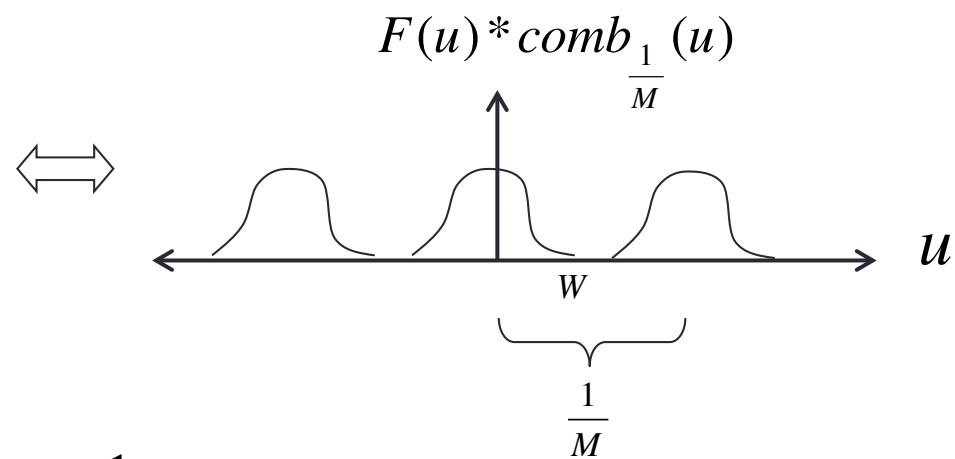
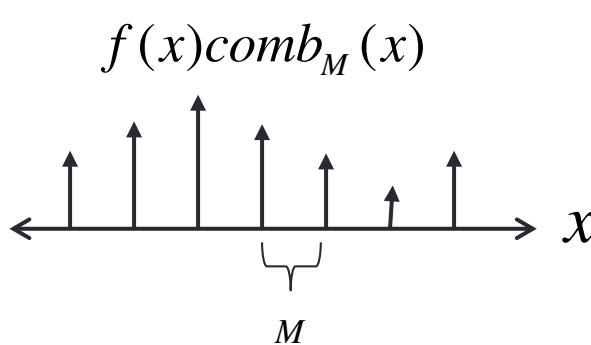
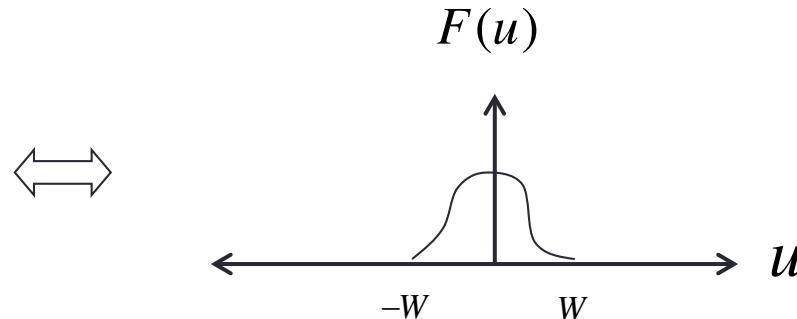
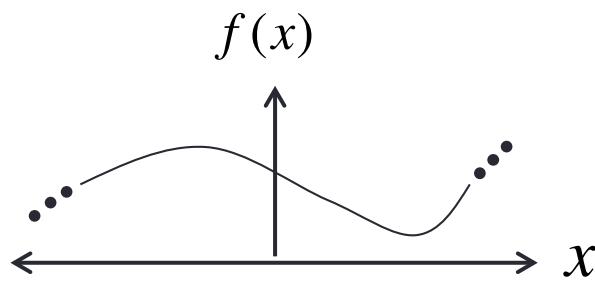
Multiply:



Convolve:

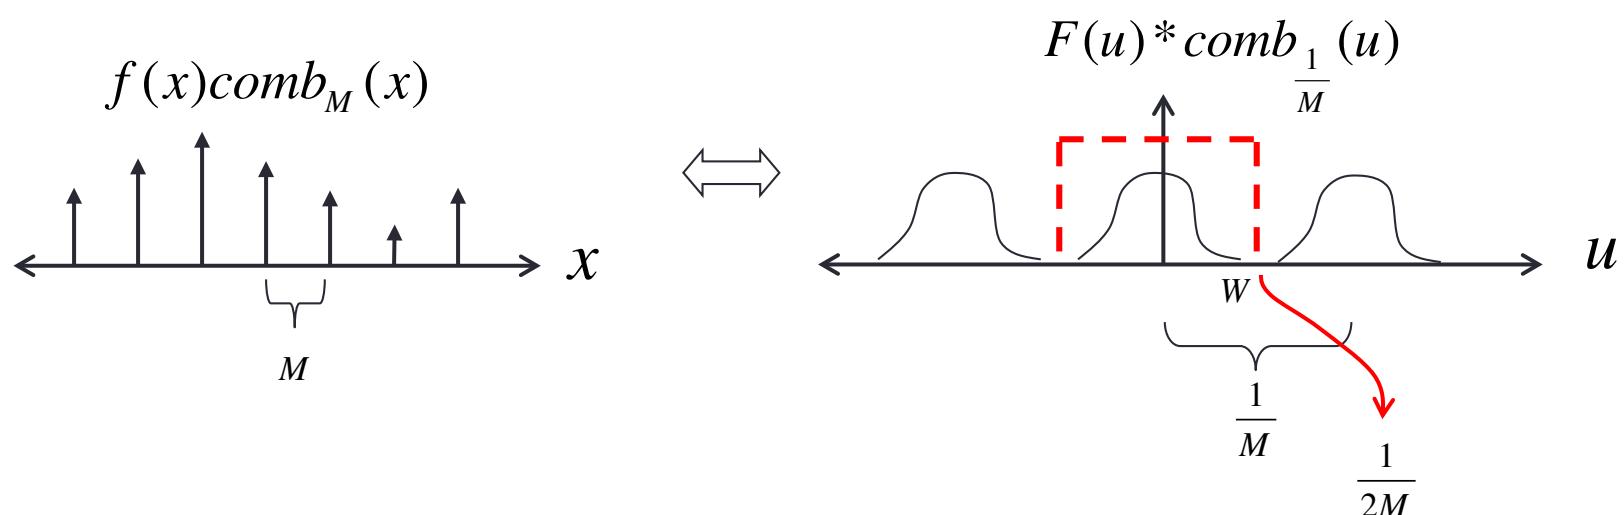


# Sampling low frequency signal



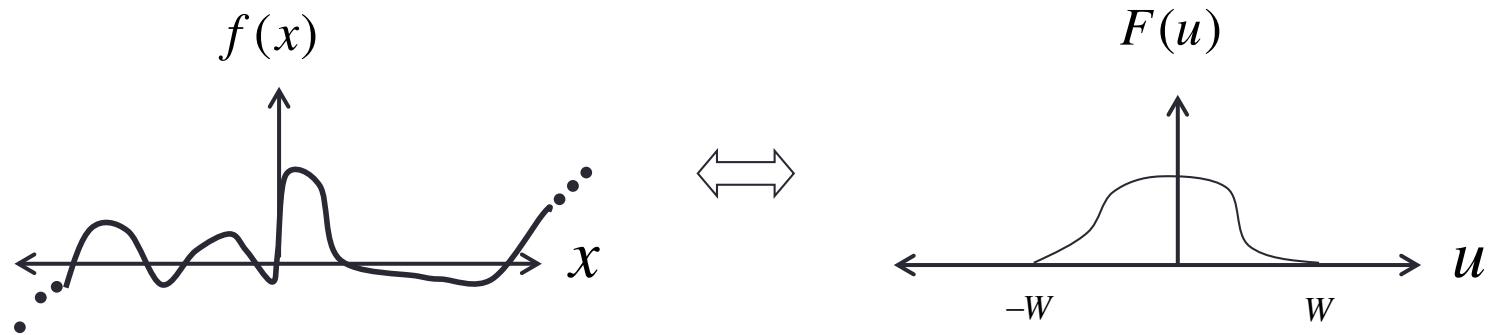
No “problem” if  $\frac{1}{M} > 2W$

# Sampling low frequency signal

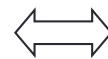


*If there is no overlap, the original signal can be recovered from its samples by low-pass filtering.*

# Sampling high frequency signal



$$f(x)comb_M(x)$$



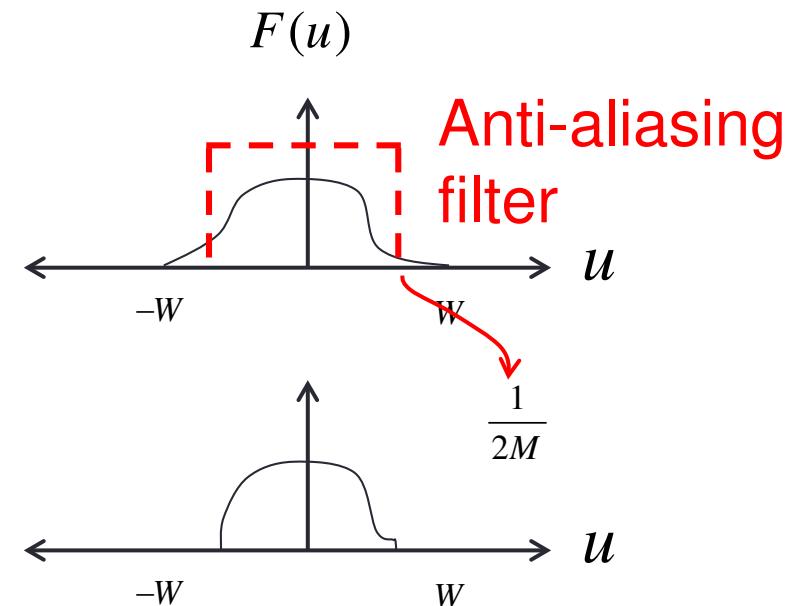
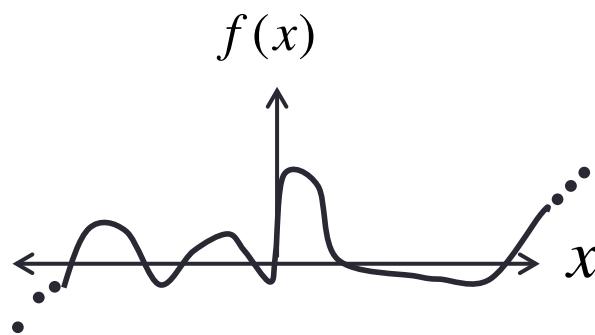
$$F(u) * comb_{\frac{1}{M}}(u)$$



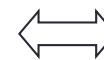
$$\frac{1}{M}$$

Overlap: The high frequency energy is folded over into low frequency. It is “aliasing” as lower frequency energy. And you cannot fix it once it has happened.

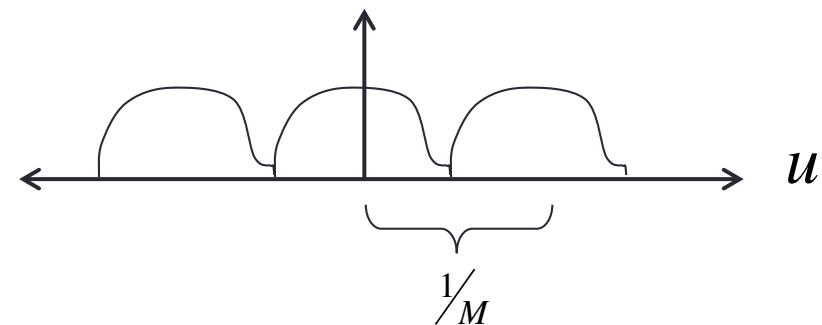
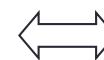
# Sampling high frequency signal



$$f(x) * h(x)$$



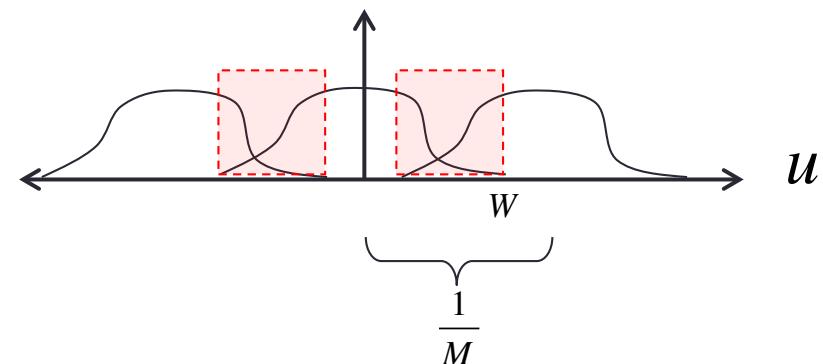
$$[f(x) * h(x)] \text{comb}_M(x)$$



# Sampling high frequency signal

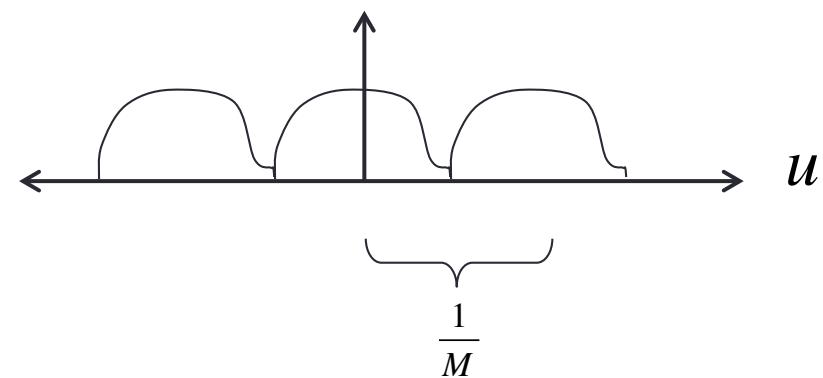
- Without anti-aliasing filter:

$$f(x) \text{comb}_M(x)$$



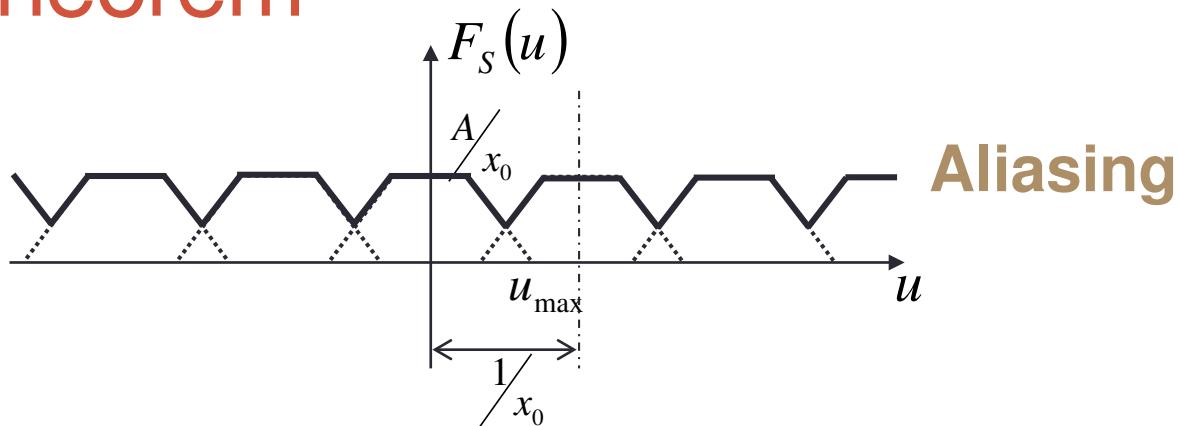
- With anti-aliasing filter:

$$[f(x) * h(x)] \text{comb}_M(x)$$



# Nyquist Theorem

If  $u_{\max} > \frac{1}{2x_0}$



When can we recover  $F(u)$  from  $F_S(u)$  ?

Only if  $u_{\max} \leq \frac{1}{2x_0}$  (Nyquist Frequency)

We can use

$$C(u) = \begin{cases} x_0 & |u| < \frac{1}{2x_0} \\ 0 & \text{otherwise} \end{cases}$$

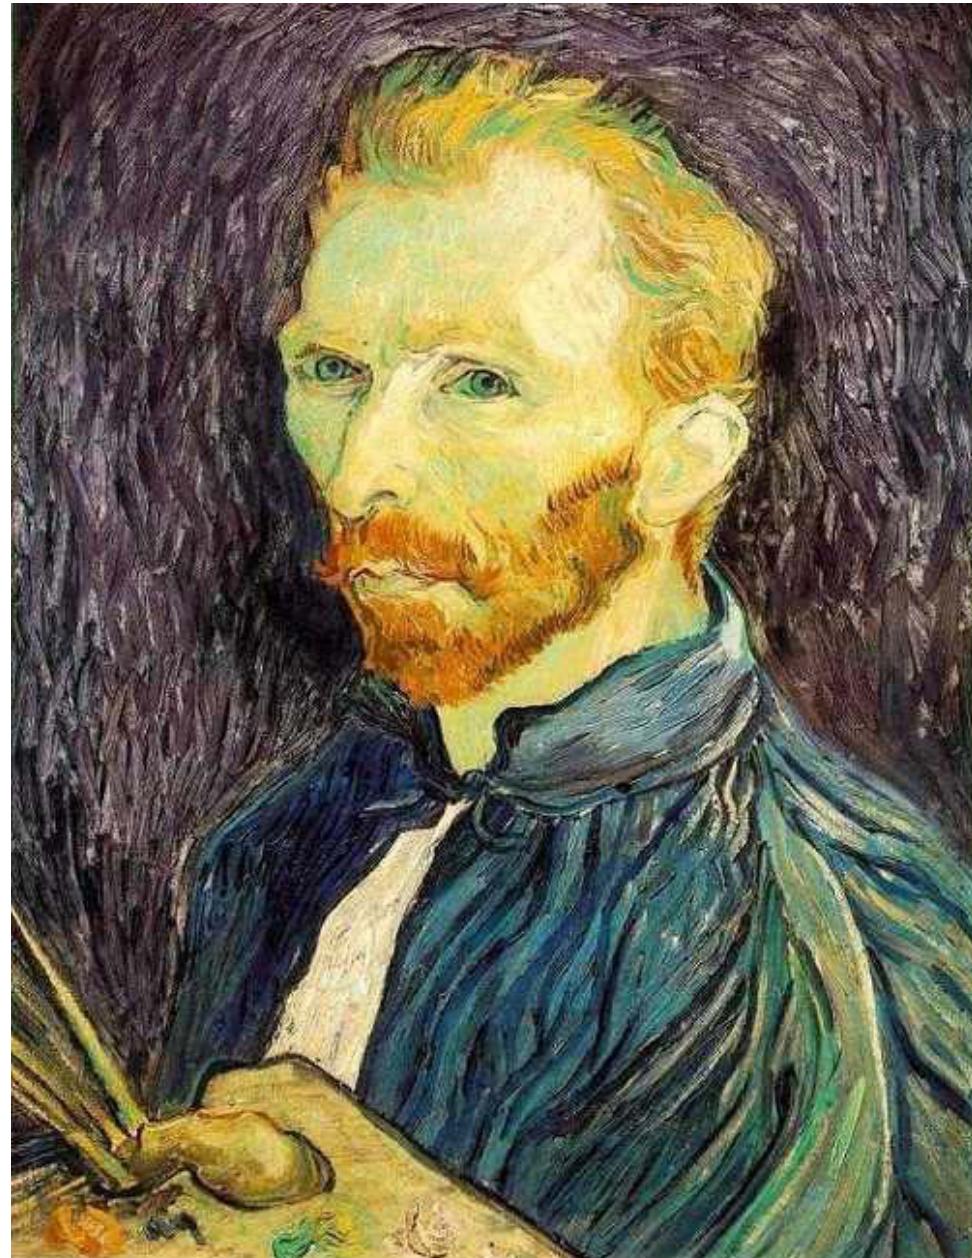
Then  $F(u) = F_S(u)C(u)$  and  $f(x) = \text{IFT}[F(u)]$

Sampling frequency must be greater than  $2u_{\max}$

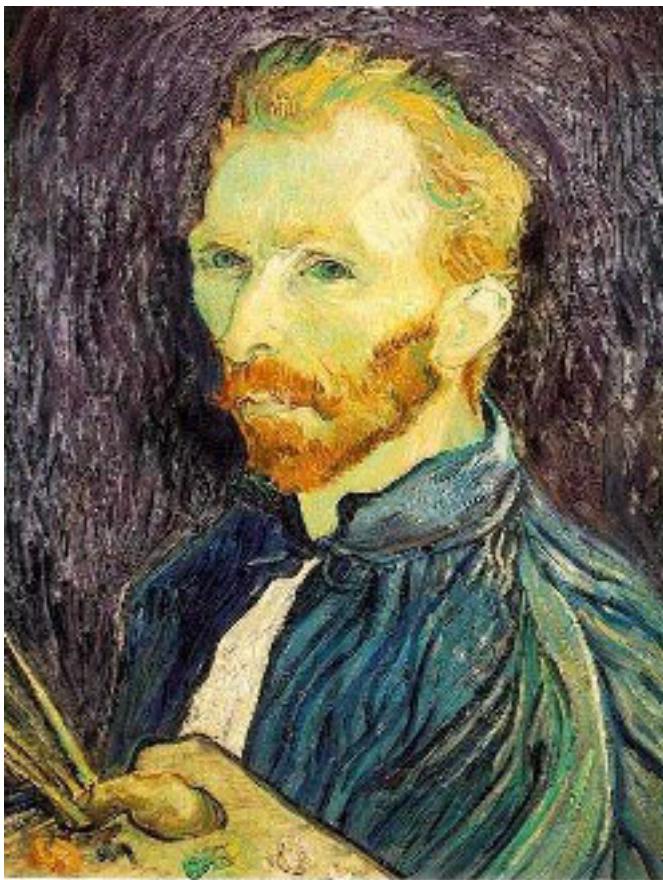
# Image half-sizing

This image is too big to fit on the screen. How can we reduce it?

How to generate a half-sized version?



# Image sub-sampling



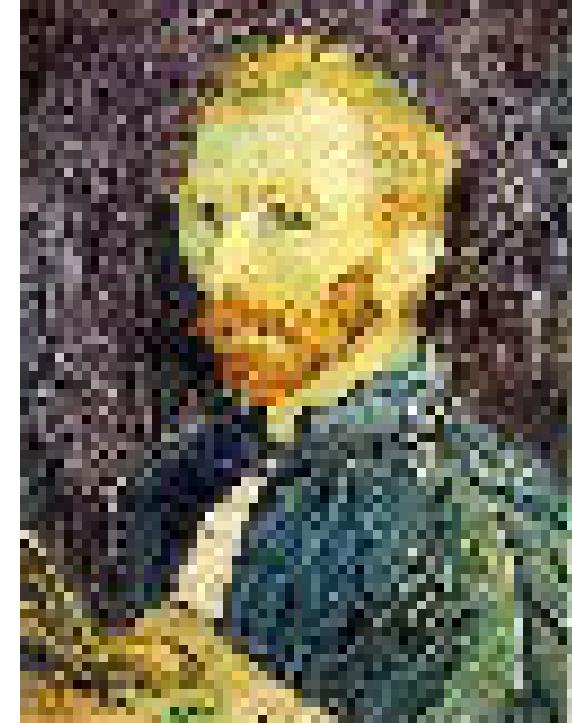
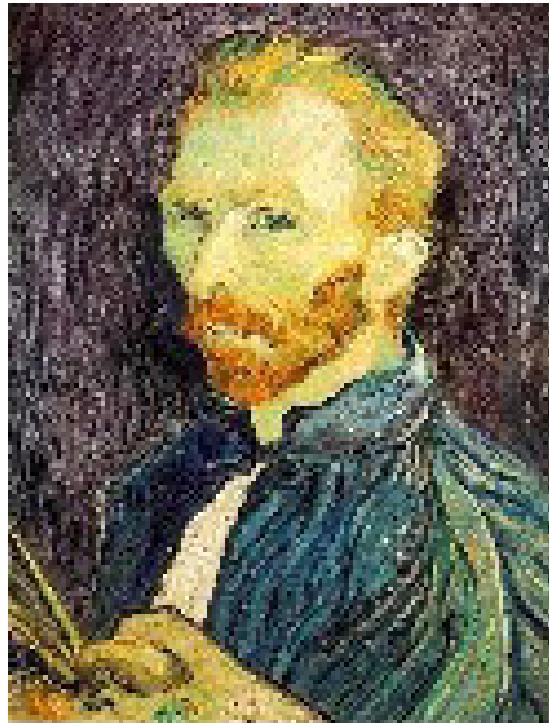
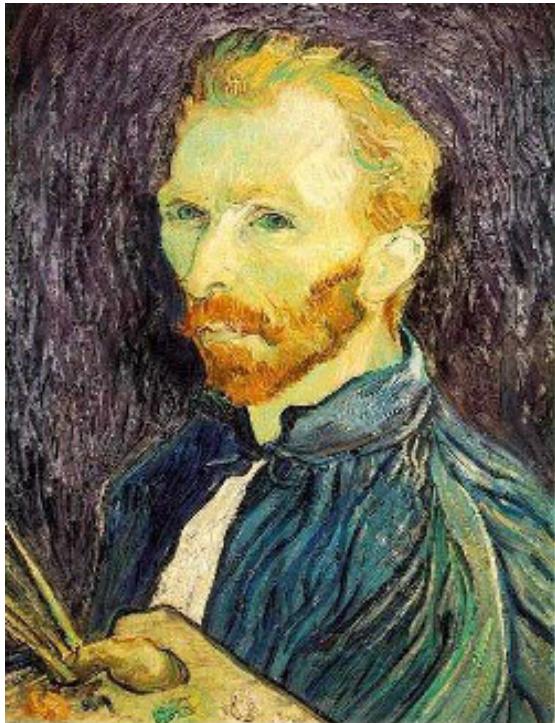
1/2



1/4

Throw away every other row and column to create a 1/2 size image  
- called *image sub-sampling*

# Image sub-sampling



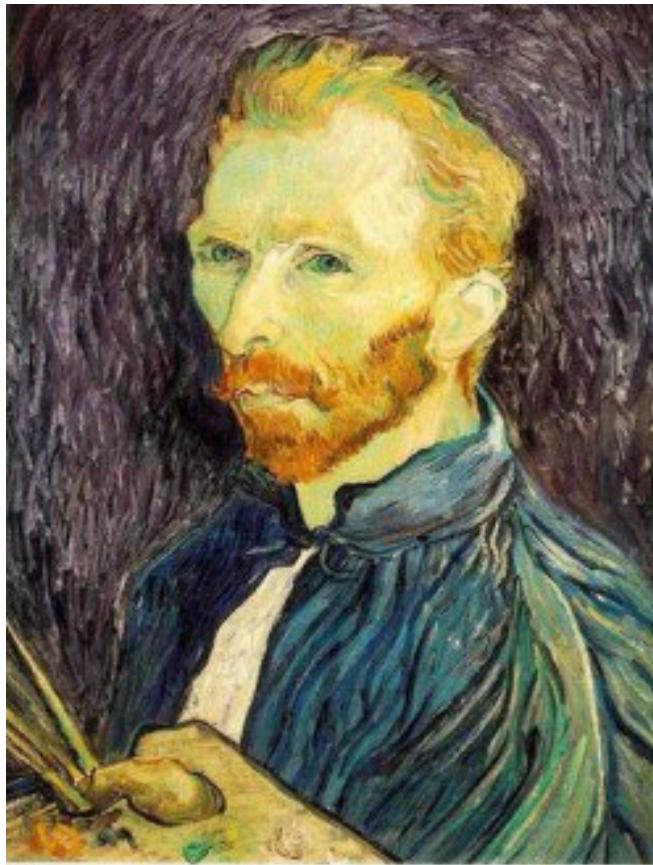
1/2

1/4 (2x zoom)

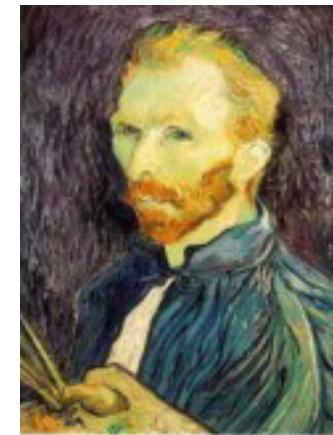
1/8 (4x zoom)

Aliasing! What do we do?

# Gaussian (lowpass) pre-filtering



Gaussian 1/2



G 1/4

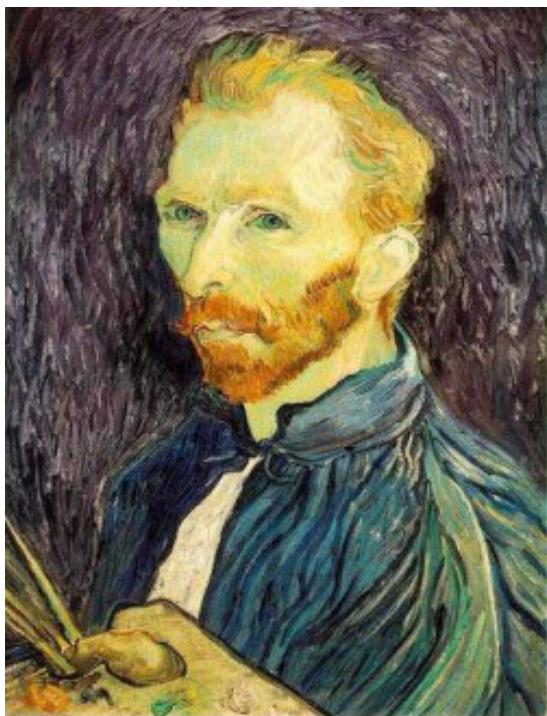


G 1/8

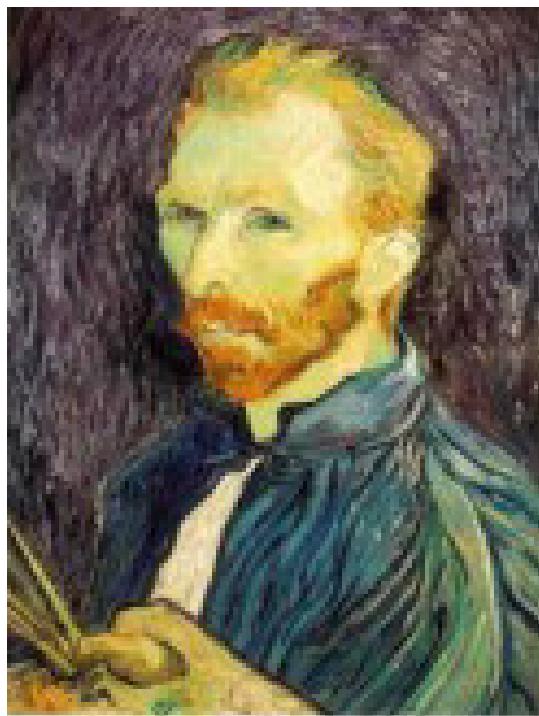
Solution: filter the image, *then* subsample

- Filter size should double for each  $\frac{1}{2}$  size reduction. Why?

# Subsampling with Gaussian pre-filtering



Gaussian 1/2

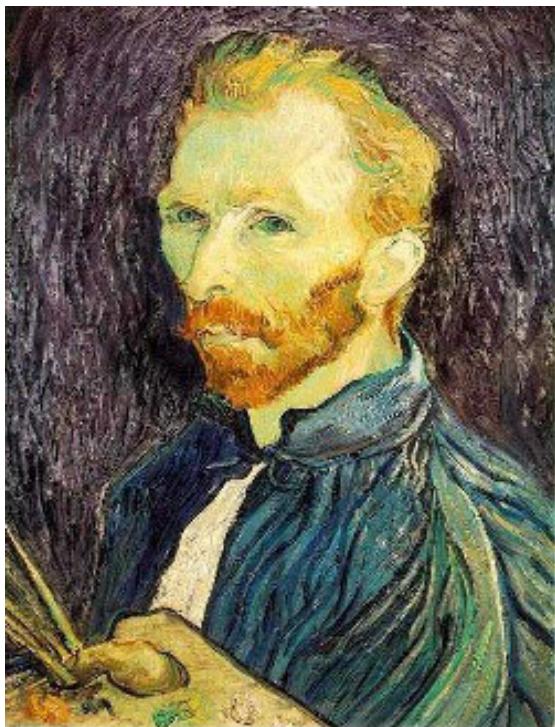


G 1/4



G 1/8

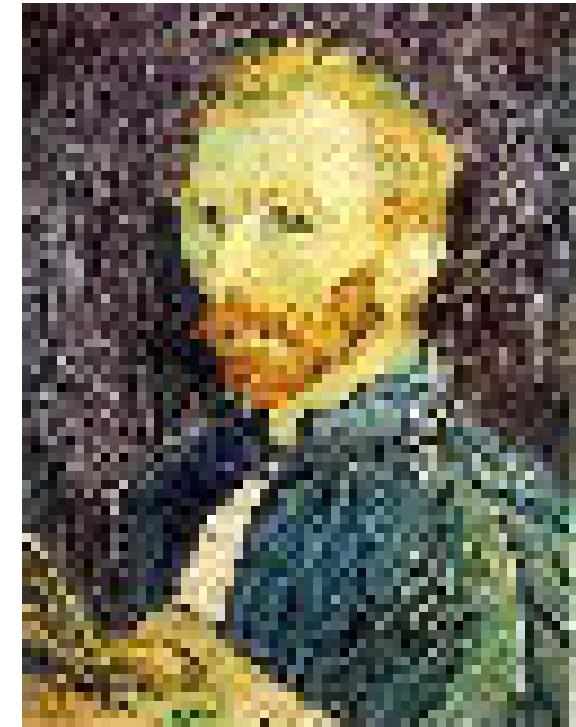
# Compare with...



1/2

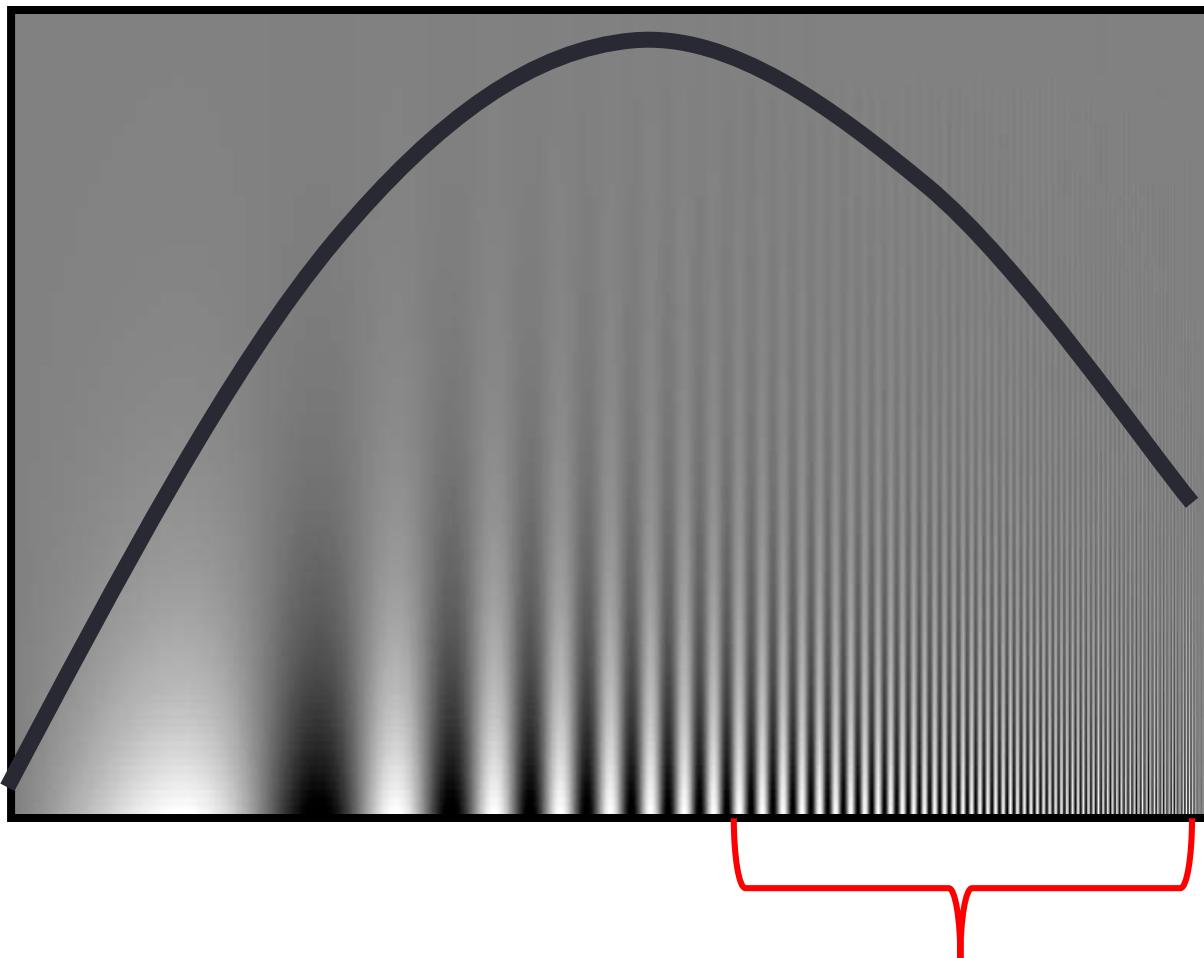


1/4 (2x zoom)



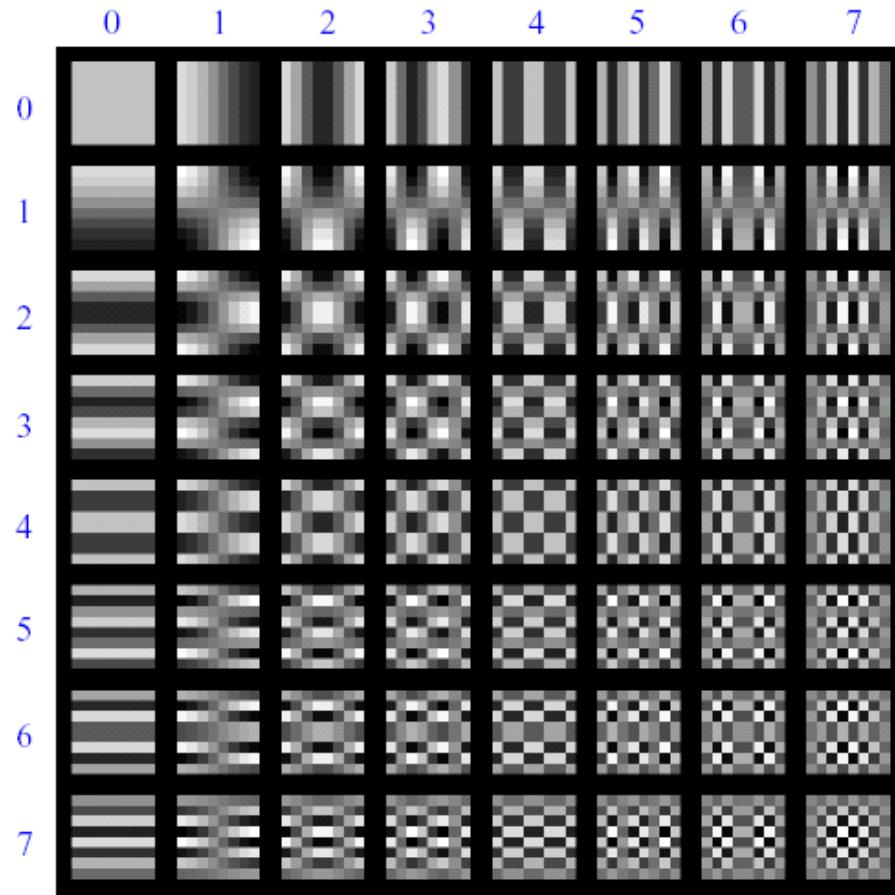
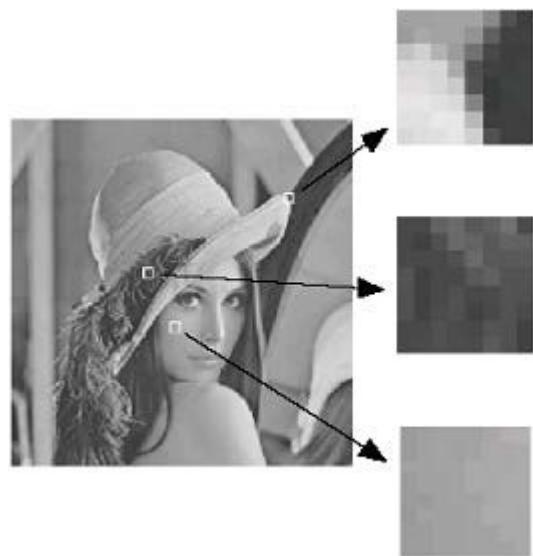
1/8 (4x zoom)

# Campbell-Robson contrast sensitivity curve



*The higher the frequency the less sensitive  
human visual system is...*

# Lossy Image Compression (JPEG)



Block-based Discrete Cosine Transform (DCT) on 8x8